

MongoDB Clustering/Sharding Lab

Comparing Single Instance vs. Sharded Cluster Performance on Big Data

(Paste your Screenshots and relevant logs in the very end)

Objective

Measure and compare the performance of a single MongoDB instance against a sharded cluster using a 100,000 document dataset.

Verify your setup:

```
docker --version
```

```
docker-compose --version
```

Lab files structure:

```
mongodb-lab/
```

```
    ├── docker-compose.single.yml
    ├── docker-compose.cluster.yml
    └── scripts/
        ├── init_single.py
        └── init_cluster.py
```

Ensure docker-compose files are in root and Python scripts are in the `scripts/` folder.

Architecture Overview

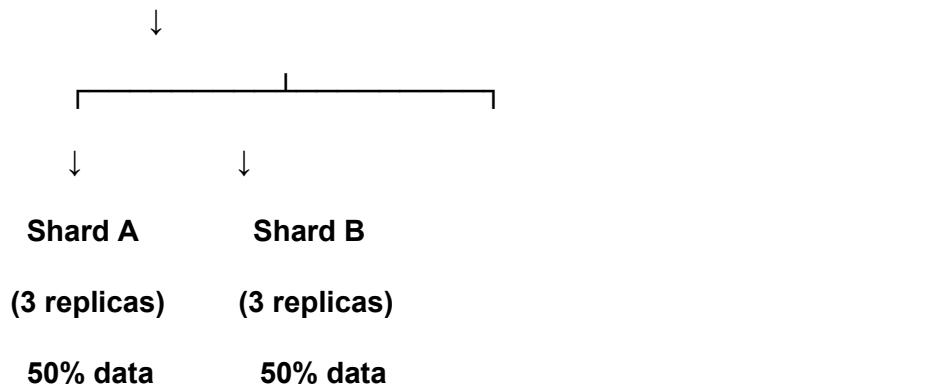
Single Instance Setup:

Application → MongoDB → All Data

Drawback: Single point of failure, limited by one machine's resources.

Sharded Cluster Setup:

Application → Mongos Router



Key components:

- **Mongos:** Query router directing traffic to appropriate shards
- **Config Servers:** Metadata storage (3-node replica set)
- **Shards:** Data storage nodes (3-node replica sets)
- **Shard Key:** Distribution field (userId, hashed)

Experiment 1: Single Instance Baseline

Run a single MongoDB server and benchmark insert, read, update, and aggregation operations.

```
# Start single instance
```

```
docker-compose -f docker-compose.single.yml up
```

```
# View output (wait 3-5 minutes for completion)
```

```
# Check results
```

```
docker logs init-single
```

```
# Save for comparison
```

```
docker logs init-single > single_results.txt
```

```
# Stop and cleanup
```

```
docker-compose -f docker-compose.single.yml down -v
```

Record these metrics:

- Insert throughput (ops/sec)
- Read throughput (queries/sec)
- Update throughput (ops/sec)

- Aggregation speed (docs/sec)

Experiment 2: Sharded Cluster

Deploy a 10-container cluster (3 config servers, 6 shard nodes, 1 router, 1 initialization script).

Start cluster

```
docker-compose -f docker-compose.cluster.yml up
```

The initialization process:

- Sets up config server replica set

- Initializes shard A (3 nodes)

- Initializes shard B (3 nodes)

- Registers shards with cluster

- Enables sharding on database

- Runs performance benchmark

View results

```
docker logs init-cluster
```

Save results

```
docker logs init-cluster > cluster_results.txt
```

Cleanup

```
docker-compose -f docker-compose.cluster.yml down -v
```

Note the following:

- Replica set initialization sequence
- Data split between shards (target: 50/50)
- Performance differences from single instance

Cluster Inspection (Required)

Start cluster in detached mode

```
docker-compose -f docker-compose.cluster.yml up -d
```

```
sleep 60
```

View cluster configuration

```
docker exec -it mongos mongo --eval "sh.status()"
```

Check shard A status

```
docker exec -it shardA1 mongo --port 27018 --eval "rs.status()"
```

Check shard B status

```
docker exec -it shardB1 mongo --port 27018 --eval "rs.status()"
```

Interactive shell

```
docker exec -it mongos mongo
```

Inside mongo shell:

```
use analyticsDB
```

```
db.events.count()
```

```
db.events.getShardDistribution()  
db.events.find({userId: 42}).limit(5).pretty()  
exit
```

OR execute commands directly from terminal:

```
docker exec -it mongos mongo analyticsDB --eval "db.events.count()"  
docker exec -it mongos mongo analyticsDB --eval "db.events.getShardDistribution()"  
docker exec -it mongos mongo analyticsDB --eval "db.events.find({userId: 42}).limit(5).pretty()"
```

Command reference:

- `sh.status()` - Cluster configuration and shard info
- `rs.status()` - Replica set member status (PRIMARY/SECONDARY roles)
- `getShardDistribution()` - Document distribution across shards
- `count()` - Total document count

Results Analysis : Fill this up and submit along side the exercises

Fill in your measurements:

Operation	Single Instance	Cluster	Improvement %
Insert (ops/sec)			%
Read (qps)			%
Update (ops/sec)			%
Aggregation (docs/sec)			%

Extract data from logs:

```
grep "ops/sec|qps|docs/sec" single_results.txt
```

```
grep "ops/sec|qps|docs/sec" cluster_results.txt
```

Typical improvements observed:

- Insert: 40-60% (parallel disk writes)
- Read: 30-40% (distributed load)
- Update: 40-60% (reduced lock contention)
- Aggregation: 40-60% (parallel computation)

Check data distribution:

```
grep "DATA DISTRIBUTION" cluster_results.txt -A 5
```

Expected result: approximately 50,000 documents per shard.

Performance Factors

1. Parallel Write Operations

- Single: 100K inserts to one disk
- Cluster: 50K to shard A + 50K to shard B concurrently

2. Read Load Distribution

- Single: All queries hit one server
- Cluster: Queries spread across 6 nodes

3. Write Lock Management

- Single: All updates queue for single write lock
- Cluster: Updates distributed, reducing contention per shard

4. Aggregation Pipeline

- Single: One CPU processes 100K documents
- Cluster: Each shard processes 50K in parallel, mongos merges results

Hashed Shard Key Implementation:

```
key={'userId': 'hashed'}
```

The hash function distributes documents evenly regardless of userId values. This prevents scenarios where sequential IDs would all land on the same shard.

Replica Set Behavior:

- Each shard: 1 PRIMARY (accepts writes) + 2 SECONDARY (replicas)
- Automatic failover if PRIMARY becomes unavailable
- Election time: typically under 30 seconds

Cleanup Procedure

```
docker-compose -f docker-compose.single.yml down -v  
docker-compose -f docker-compose.cluster.yml down -v  
docker ps -a | grep mongo  
docker rm -f $(docker ps -aq --filter "name=mongo")  
docker system prune -f
```

Lab Questions

Q1: Performance Analysis Which operation showed the greatest performance gain in the cluster? Explain why based on the architecture differences.

Q2: Data Distribution What problems would occur if shard A contained 90% of documents instead of 50%?

Q3: Shard Key Design Analyze these alternative shard key choices:

- Timestamp field
- Boolean field
- Auto-incrementing integer

What issues would each introduce?

Q4: Scaling Strategy The dataset will grow from 100K to 10M documents next year, with 10x traffic increase. Propose a scaling plan.

Q5: Failover Test Execute this failover scenario:

```
docker-compose -f docker-compose.cluster.yml up -d  
sleep 60  
docker exec -it shardA1 mongo --port 27018 --eval "rs.status()"  
docker kill shardA1  
sleep 30  
docker exec -it shardA2 mongo --port 27018 --eval "rs.status()"
```

```
docker exec -it mongos mongo --eval "db.getSiblingDB('analyticsDB').events.count()"  
docker start shardA1
```

Measure the downtime and describe shardA1's role after restart.

Q6: Query Routing Compare these queries using `.explain("executionStats")`:

You can directly run from terminal too:

```
db.events.find({userId: 42})      // targeted  
db.events.find({eventType: "click"}) // scatter-gather
```

How many shards does each query hit?

Q7: Bottleneck Analysis Identify the primary bottleneck in the single instance (CPU/disk/RAM/locks). Cite specific log evidence.

Q8: Infrastructure Economics The cluster uses 10x the containers of single instance. In what scenarios does this cost multiply justify itself?

Command Reference

Single instance workflow

```
docker-compose -f docker-compose.single.yml up
```

```
docker logs init-single
```

```
docker-compose -f docker-compose.single.yml down -v
```

Cluster workflow

```
docker-compose -f docker-compose.cluster.yml up
```

```
docker logs init-cluster
```

```
docker exec -it mongos mongo --eval "sh.status()"
```

```
docker-compose -f docker-compose.cluster.yml down -v
```

System cleanup

```
docker system prune -a -f
```

End of Lab

```
protocol": "op_msg", "durationMillis": 139}]  
init-single | ======  
init-single | SINGLE MONGODB INSTANCE - PERFORMANCE BASELINE  
init-single | ======  
init-single | Architecture: Single mongod process  
init-single | URI: mongodb://mongo-single:27017  
init-single | Dataset: 100,000 documents (simulating big data workload)  
init-single | ======  
init-single | ✓ Connected to MongoDB  
init-single | ✓ Collection cleaned  
init-single | [1/4] INSERT TEST: 100,000 documents  
-----  
init-single | Simulating: Real-time event streaming (clicks, purchases, views)  
init-single | Progress: 20,000 docs | Current rate: 85,047 ops/sec  
init-single | Progress: 40,000 docs | Current rate: 90,668 ops/sec  
init-single | Progress: 60,000 docs | Current rate: 88,747 ops/sec  
init-single | Progress: 80,000 docs | Current rate: 89,601 ops/sec  
init-single | Progress: 100,000 docs | Current rate: 88,892 ops/sec  
-----  
init-single | ✓ INSERT COMPLETED  
init-single | Time taken: 1.12 seconds  
init-single | Throughput: 88,890 ops/sec  
init-single | Total documents: 100,000  
init-single | Bottleneck: Single node disk I/O + CPU  
-----  
init-single | [2/4] READ TEST: 10,000 targeted queries (by userId)  
-----  
init-single | Simulating: User activity lookups, analytics dashboards  
init-single | Progress: 2,000 queries | Current QPS: 242  
init-single | Progress: 4,000 queries | Current QPS: 235  
init-single | Progress: 6,000 queries | Current QPS: 237
```

✓ READ COMPLETED

Time taken: 39.63 seconds
Query rate: 252 queries/sec
Avg latency: 3.96 ms/query
Advantage: Load balanced across replica sets

[3/4] UPDATE TEST: 10,000 random updates

Benefit: Updates distributed to respective shard primaries

Progress: 2,000 updates | Current rate: 1,758 ops/sec
Progress: 4,000 updates | Current rate: 1,684 ops/sec
Progress: 6,000 updates | Current rate: 1,585 ops/sec
Progress: 8,000 updates | Current rate: 1,510 ops/sec
Progress: 10,000 updates | Current rate: 1,473 ops/sec

✓ UPDATE COMPLETED

Time taken: 6.79 seconds
Update rate: 1,473 ops/sec
Advantage: No single-node write lock contention

[4/4] AGGREGATION TEST: Complex analytics query

Benefit: Parallel processing across all shards

Computed: Event type x Device breakdown
Records processed: 100,000 (across 2 shards)
Results returned: 15

✓ AGGREGATION COMPLETED

Time taken: 0.14 seconds
Processing rate: 729,234 docs/sec
Advantage: Parallel aggregation pipeline on each shard

```
init-cluster |  
init-cluster | ======  
init-cluster | MONGODB SHARDED CLUSTER - SETUP AND PERFORMANCE TEST  
init-cluster | ======  
init-cluster |  
init-cluster | [1/5] Initializing Config Server Replica Set (3 nodes)  
init-cluster | -----  
init-cluster |   Connecting to mongodb://cfg1:27019...  
init-cluster |     ✓ Initiated replica set: cfgRepl  
init-cluster |     ✓ Primary elected in replica set  
init-cluster |  
init-cluster | [2/5] Initializing Shard A Replica Set (3 nodes)  
init-cluster | -----  
init-cluster |   Connecting to mongodb://shardA1:27018...  
init-cluster |     ✓ Initiated replica set: shardA  
init-cluster |     ✓ Primary elected in replica set  
init-cluster |  
init-cluster | [3/5] Initializing Shard B Replica Set (3 nodes)  
init-cluster | -----  
init-cluster |   Connecting to mongodb://shardB1:27018...  
init-cluster |     ✓ Initiated replica set: shardB  
init-cluster |     ✓ Primary elected in replica set  
init-cluster |  
init-cluster | [4/5] Adding Shards to Cluster via Mongos Router  
init-cluster | -----  
init-cluster |   Connecting to mongos: mongodb://mongos:27017  
init-cluster |     ✓ Added shardA to cluster  
init-cluster |     ✓ Added shardB to cluster  
init-cluster |  
init-cluster | [5/5] Configuring Database Sharding  
init-cluster | -----  
init-cluster |     ✓ Enabled sharding on database: analyticsDB  
init-cluster |     ✓ Sharded collection on userId (hashed)  
init-cluster |  
init-cluster |     ✓ Cluster setup complete!  
init-cluster |  
init-cluster | ======  
init-cluster | SHARDED CLUSTER - PERFORMANCE TEST  
init-cluster | ======
```

[3/4] UPDATE TEST: 10,000 updates

```
Progress: 2,000 updates | Rate: 589 ops/sec
Progress: 4,000 updates | Rate: 579 ops/sec
Progress: 6,000 updates | Rate: 581 ops/sec
Progress: 8,000 updates | Rate: 577 ops/sec
Progress: 10,000 updates | Rate: 578 ops/sec
```

✓ UPDATE COMPLETED

Time: 17.31s | Throughput: 578 ops/sec

[4/4] AGGREGATION TEST

```
Processed: 100,000 docs
Results: 15 groups
```

✓ AGGREGATION COMPLETED

Time: 0.13s | Rate: 748,066 docs/sec

=====

CLUSTER RESULTS SUMMARY

=====

```
Insert: 56,252 ops/sec
Read: 231 queries/sec
Update: 578 ops/sec
Aggregation: 748,066 docs/sec
```

=====

```
MongoDB server version: 4.4.29
--- Sharding Status ---
sharding version: {
    "_id" : 1,
    "minCompatibleVersion" : 5,
    "currentVersion" : 6,
    "clusterId" : ObjectId("6922b0a5a5e7a5c8a000fa56")
}
shards:
{ "_id" : "shardA", "host" : "shardA/shardA1:27018,shardA2:27018,shardA3:27018", "state" :
{ "_id" : "shardB", "host" : "shardB/shardB1:27018,shardB2:27018,shardB3:27018", "state" :
active mongoses:
    "4.4.29" : 1
autosplit:
    Currently enabled: yes
balancer:
    Currently enabled: yes
    Currently running: no
    Failed balancer rounds in last 5 attempts: 0
    Migration Results for the last 24 hours:
        512 : Success
databases:
{ "_id" : "analyticsDB", "primary" : "shardA", "partitioned" : true, "version" : { "uuid" :
: Timestamp(1763881152, 5) }
{ "_id" : "config", "primary" : "config", "partitioned" : true }
    config.system.sessions
        shard key: { "_id" : 1 }
        unique: false
        balancing: true
        chunks:
            shardA 512
            shardB 512
        too many chunks to print, use verbose if you want to force print
MongoDB shell version v4.4.29
connecting to: mongodb://127.0.0.1:27018/?compressors=disabled&gssapiServiceName=mongodb
Implicit session: session { "id" : UUID("c16c1db7-05b6-4d5a-a106-cbf862148c4c") }
MongoDB server version: 4.4.29
{
    "set" : "shardA",
    "date" : ISODate("2025-11-23T07:26:47.293Z")
```