

Big Data Analytics

Fall 2025

Lecture 8

Dr. Tariq Mahmood



Hadoop

- When a dataset outgrows the storage capacity of a single physical machine, it becomes necessary to partition it across a number of separate machines.
- Filesystems that manage the storage across a network of machines are called **distributed filesystems**.
- Since they are network based, **all the complications of network programming kick in**, thus making distributed filesystems more complex than regular disk filesystems.
- Making the filesystem tolerate to node failure without suffering data loss.



HDFS

- HDFS: Hadoop Distributed Filesystem.
- HDFS: filesystem designed for storing very large files with streaming data access patterns, running on clusters of commodity hardware
- Very large files: hundreds of MB, GB, TB or PB in size
- Streaming data access
 - Write-once, read-many-times pattern.
 - A dataset is copied from source and analyzed
 - Each analysis will involve a large proportion, if not all, of the dataset, so the time to read the whole dataset is more important than the latency in reading the first record.



HDFS

- **Commodity hardware**
 - Doesn't require expensive, highly reliable hardware.
 - Designed to run on clusters of commodity hardware for which the chance of node failure across the cluster is high, at least for large clusters.
 - HDFS is designed to carry on working without a noticeable interruption
- **Not a good-fit for Low-latency data access:**
 - Applications that require low-latency access (milliseconds range), will not work with HDFS
 - HDFS: optimized for delivering high throughputs at the expense of latency




HDFS

- Not a good-fit for lots of small files:
 - Because the namenode holds HDFS metadata in memory - the limit to the number of files is governed by the amnt of memory on namenode.
 - ROT: Each HDFS block's metadata takes about 150 bytes on namenode
 - So, for example, if you had one million files, each taking one block, you would need at least 150 MB of memory.
 - $150 * 1 \text{ million} = 150 \text{ million}$
 - $1 \text{ MB} = 1048576 \text{ bytes}$
 - So mem = 143.05 MB



Blocks

- A disk has a **block** size, which is the minimum amount of data that it can read or write.
- Filesystems also deal with data in blocks, which are an **integral multiple of the disk block size**.
- Filesystem blocks: typically a few kilobytes in size, whereas disk blocks are normally 512 bytes.




□ The thicker lines divides the sectors



Blocks

- HDFS: concept of a block, but much larger unit—128 MB by default.
- Files in HDFS are broken into block-sized chunks - stored as independent units.
- Unlike a filesystem for a single disk, **a file in HDFS that is smaller than a single block does not occupy a full block's worth of underlying storage.**
- For example, a 1 MB file stored with a block size of 128 MB uses 1 MB of disk space, not 128 MB.





Why Is a Block in HDFS So Large?

- **To minimize the cost of seeks:** If the block is large, the time it takes to transfer the data from the disk can be longer than the time to seek to the start of the block.
- Transferring a large file made of multiple blocks operates at the disk transfer rate.
- If seek time is around 10 ms, and the transfer rate is 100 MB/s, to make the seek time 1% of the transfer time, we need to make the block size around 100 MB: default is actually 128 MB



Why Is a Block in HDFS So Large?

$$T_{seek} = 0.01 \text{ s}$$

$$T_{transfer} = \frac{\text{Block Size}}{\text{Transfer Rate}}$$

$$T_{seek} = 0.01 \times T_{transfer}$$

| File Size | Configured Block Size | No. of Blocks | Block Sizes |
|-----------|-----------------------|---------------|--------------|
| <120 MB | 128 MB | 1 | 120 MB |
| =128 MB | 128 MB | 1 | 128 MB |
| =134 MB | 128 MB | 2 | 128 MB, 6 MB |
| =512 MB | 128 MB | 4 | 128 MB × 4 |



Benefits

- **First:** a file can be larger than any single disk in the network.
 - There's nothing that requires the blocks from a file to be stored on the same disk, so they can take advantage of any of the disks in the cluster.
- **Second:** making the unit of abstraction a block rather than a file simplifies the storage subsystem.
- **Third:** blocks fit well with replication for providing fault tolerance and availability.



Benefits

- Like its disk filesystem cousin, HDFS's fsck command understands blocks.
- % hdfs fsck / -files -blocks
- This will list the blocks that make up each file in the filesystem

```
[manjunath@engoid-sas-stage-host-005 ~]$ hdfs fsck /tmp/data
Connecting to namenode via http://[REDACTED]:fsck?ugi=manjunath
FSCK started by manjunath (auth:SIMPLE) from /192.0.1.9 for path /tmp/data at
.Status: HEALTHY
Total size: 15882 B
Total dirs: 8
Total files: 1
Total symlinks: 0
Total blocks (validated): 1 (avg. block size 15882 B)
Minimally replicated blocks: 1 (100.0 %)
Over-replicated blocks: 0 (0.0 %)
Under-replicated blocks: 0 (0.0 %)
Mis-replicated blocks: 0 (0.0 %)
Default replication factor: 1
Average block replication: 1.0
Corrupt blocks: 0
Missing replicas: 0 (0.0 %)
Number of data-nodes: 3
Number of racks: 1
FSCK ended at Mon Jun 28 08:12:11 UTC 2021 in 0 milliseconds

The filesystem under path '/tmp/data' is HEALTHY
[manjunath@engoid-sas-stage-host-005 ~]$
```



Questions

- Scenario: You have a file of size 800 MB, and the HDFS block size is set to 128 MB.
- Question: How many blocks will this file be split into on HDFS?
- Number of blocks= $800/128 = 6.25$
- HDFS rounds this up, so the file will be split into **7 blocks**: 6 full blocks of 128 MB
- 1 block of 16 MB (remainder)

Questions

- Scenario: You have a file of size 1 GB stored in HDFS, and the replication factor is set to 3.
- Question: How much disk space will be used to store this file?
- 3 GB

Questions

- Scenario: You have 5 files, each of size 500 MB, and the HDFS block size is set to 256 MB.
- Question: How many blocks are needed to store all the files?
- $500/256 \sim 2$ blocks per file (in reality, 1 block + 0.953125 blocks)
- For 5 files ~ 10 blocks

Questions

- Scenario: You have a file of size 10 MB, and the HDFS block size is 128 MB. The replication factor is set to 2.
- Question: How much disk space will this file consume?
- Since the block size is 128 MB, the file will consume 1 block even though it is smaller than the block size (in this case, < 1 block is not possible)
- With a replication factor of 2: 256MB


Questions

- Scenario: You have a file of size 5 GB, and the HDFS block size is 512 MB.
- Question: How many blocks will this file occupy in HDFS?

- $5000 \text{ MB} / 512 \text{ MB} = 9.77 \sim 10 \text{ blocks}$
- 9 full blocks of 512 MB
- 1 block of 232 MB

Questions

- Scenario: You have a file of size 2.5 GB, and HDFS supports two block sizes: 128 MB and 256 MB.
- Question: How many blocks are required to store the file with each block size?
- Scenario: You have 10 files, each of size 300 MB, the HDFS block size is set to 128 MB, and the replication factor is set to 3.
- Question: How much total disk space will be consumed to store these files in HDFS?



| NameNode Task | Task Details | Explanation |
|--|--|---|
| 1. Namespace Management | Maintains the entire HDFS directory tree — file names, directories, permissions, timestamps, and ownership. | Think of it as the “file system table” (like ls output for the whole cluster). |
| 2. Block Management | Keeps a mapping of each file to its blocks and each block to its DataNodes. | If a file is 1 GB and split into 8 blocks, the NameNode knows which DataNodes hold each of those 8 blocks. |
| 3. Replication Control | Ensures every block has the configured replication factor (default = 3). If a DataNode fails, it schedules replication from other nodes to restore balance. | Like a “health monitor” maintaining 3 copies of each piece. |
| 4. Metadata Persistence (EditLog + FsImage) | Periodically writes metadata to disk: FsImage: snapshot of the namespace EditLog: record of recent changes It | Merges them periodically into a new FsImage (checkpointing). Ensures HDFS can recover state after restart or crash. |
| 5. Client Coordination & Access Control | When a client requests to read/write a file, the NameNode: Authenticates the user Returns block locations for reading Chooses target DataNodes for new block writes | Acts as a “traffic controller” that coordinates clients and DataNodes without transferring the actual data. |

| DataNode Tasks | Task Details | Explanation |
|--|--|--|
| 1. Block Storage and Retrieval | Stores the actual HDFS blocks (e.g., 128 MB each) on local disks and reads/writes them upon NameNode or client request. | Acts like the “warehouse shelves” that hold real data chunks. |
| 2. Block Reporting | Periodically sends a Block Report to the NameNode containing a list of all blocks it currently holds. | Like an inventory list telling the central manager (NameNode) what's in stock. |
| 3. Heartbeat Transmission | Sends regular heartbeats (every 3 seconds by default) to the NameNode to indicate that it's alive and functioning. | A “pulse check” confirming this DataNode is healthy and reachable. |
| 4. Data Replication (and Deletion) | When instructed by the NameNode, the DataNode replicates blocks to other nodes or deletes excess and obsolete copies to maintain the desired replication factor. | Think of a warehouse worker copying or removing boxes as per central orders. |
| 5. Data Transfer between DataNodes (Pipeline) | During writes, DataNodes form a replication pipeline (Node 1 → Node 2 → Node 3) to stream data efficiently. They also handle rebalancing transfers. | Like a production line where data flows through multiple nodes to form 3 copies. |

| YARN Task | Task Details | Explanation |
|---|--|--|
| 1. Resource Management | Allocates CPU, memory, and containers across the cluster for multiple jobs. The ResourceManager (RM) acts as a central scheduler. | Like an “airport control tower” deciding which plane (job) gets which runway (resource). |
| 2. Job Scheduling & Queuing | Determines which application runs first and how cluster resources are divided among users or queues (FIFO, Capacity, Fair Scheduling). | Similar to a “task manager” prioritizing processes on your OS. |
| 3. Application Lifecycle Management | Launches, monitors, and terminates ApplicationMasters and their containers. Keeps track of app states (RUNNING, FINISHED, FAILED). | Like a project manager supervising each running job from start to finish. |
| 4. Fault Tolerance & Recovery | Detects failed ApplicationMasters or containers and re-launches them automatically on healthy nodes. Maintains high cluster uptime. | Like restarting a crashed service without user intervention. |
| 5. Resource Monitoring & Reporting | Collects usage metrics (CPU, RAM, container logs) from NodeManagers, updates cluster status, and provides data to the Web UI and job history server. | Like a system monitor showing cluster-wide performance stats. |

| NodeManager (in YARN) Tasks | Task Details | Explanation |
|---|--|--|
| 1. Container Management | Launches, monitors, and terminates containers (isolated environments where actual tasks run). Each container runs part of an application. | Like Docker running and managing containers on each node. |
| 2. Resource Monitoring | Tracks local resource usage (CPU, memory, disk, network) of all running containers and ensures no container exceeds its allocated limits. | Like a system resource monitor that enforces quotas per process. |
| 3. Node Health Reporting | Periodically sends heartbeat signals to the ResourceManager, confirming that the node is alive, and reporting container statuses. | Like sending “I’m alive and healthy” signals to the central controller. |
| 4. Log Aggregation & Cleanup | Collects logs from all containers, aggregates them, and sends them to the YARN log server (for later viewing via the Web UI). Also cleans up old containers and temporary files. | Like a janitor who collects and archives logs, then cleans up old workspace files. |
| 5. Communication with ApplicationMaster | Provides the ApplicationMaster with updates about container status, failures, and resource availability, enabling job progress tracking and retries. | Like a site supervisor updating the project manager about worker progress. |

| ApplicationMaster Tasks | Task Details | Explanation |
|--|---|---|
| 1. Resource Negotiation with ResourceManager | Communicates with the ResourceManager (RM) to request containers with specific resource requirements (CPU, memory, locality). | Like a project manager asking headquarters for workers and equipment. |
| 2. Container Launch & Coordination via NodeManagers | Once containers are allocated, the AM contacts the relevant NodeManagers to launch the tasks inside those containers. | Like assigning workers (tasks) to different construction sites (nodes). |
| 3. Task Monitoring & Progress Tracking | Tracks the status of each task running inside containers — success, failure, progress %, and time. Handles retries for failed tasks. | Like a manager checking every team member's daily progress. |
| 4. Fault Tolerance & Task Recovery | Detects failed containers or nodes and requests replacement containers from the RM to re-run failed tasks. | Like reassigning a task to another worker when one falls sick. |
| 5. Reporting & Communication with Client | Sends periodic status updates and final results back to the client program (the user who submitted the job). Updates are also written to the YARN Job | Like a manager reporting job status and results to the client. |