

Big Data Analytics

Fall 2025

Lecture 9

Dr. Tariq Mahmood



HDFS

- **WORM: Write-once, read-many-times storage pattern**
 - Once the data is recorded, it is permanently fixed. You cannot overwrite or erase it.
 - The stored data can be read, copied, or retrieved any number of times without changing it.
 - Ensures data integrity (tampering, accidental delete)



HDFS

- Commodity hardware – which are not server machines
 - Multi-core processor
 - 16-64 GB RAM
 - SSD/HDD 2 TB – but SSD for NameNode
 - 1 Gbps minimum Internet
 - Linux
- Not good for low-latency (video, Internet, online game, IoT):
 - Because designed for batch processing
- Now Hadoop with Spark/Flink is enough to handle low-latency



HDFS

- Not a good-fit for lots of small files:
 - Because the NameNode holds HDFS metadata in memory - the limit to the number of files is governed by the amnt of memory on namenode.
- ROT:
 - Each block metadata takes about 150 bytes on namenode
 - Each file inode is about 150 bytes as well
 - Total per file with 1 block = 300 bytes
- If 1 million files, each taking one block, you would need at least 300 MB of memory (calculate and verify).




Components


- An HDFS cluster: **NameNode** (the master) and a number of **datanodes** (workers).
- The NameNode **manages the filesystem (tree) namespace**.
- This is metadata which is stored as two files: **the namespace image (fsimage)** and **the edit log (edits)**
- FSimage: snapshot (checkpoint) of the entire HDFS namespace
- Edits: transaction log that records every change or modification to the HDFS namespace after the last fsimage checkpoint
- **Current state = fsimage + edit log (applied)**
- **Too much work for the NameNode!**




Components


HDFS Architecture









Components



Components




Namespace Metadata & Log





Components





Components

- A client accesses the filesystem (RPC calls) by communicating with the name-node and datanodes
- Datanodes are workhorses: store and retrieve blocks when they are told to (by clients or namenode) - report back periodically with lists of blocks that they are storing
- Without namenode, no filesystem - if namenode machine is obliterated, all files on filesystem would be lost
- Important to make the namenode resilient to failure, and Hadoop provides two mechanisms for this



Components

- Current state = fsimage + edit log (applied)
- Use a Secondary NameNode to help with checkpointing work!

Tasks
Download the current fsimage and edits from the NameNode (via HTTP).
Load them into memory and apply all edits to update the namespace.
Create a new merged fsimage.chkpoint file (a fresh checkpoint).
Send back this new fsimage to the NameNode (or store it locally).
The NameNode replaces its old fsimage and resets the edit log.




Secondary NameNode

- Periodically merge the namespace image with the edit log to prevent the edit log from becoming too large.
- Runs on a separate physical machine
- Keeps a copy of the merged namespace image (used in failure)
- State of secondary lags that of the primary – if total failure of the primary, then data loss is almost certain.
- Then: You can use local backup and transfer it to the secondary.



Components





HDFS High Availability

- The combo: replicating metadata on multiple filesystems and using secondary to create checkpoints is good - **but it does not provide high availability of the filesystem.**
- The namenode is still a **single point of failure (SPOF)**.
 - All clients—including MapReduce jobs—would be unable to read, write, or list files, because the namenode is the sole repository of the metadata and the file-to-block mapping.
- The whole Hadoop would effectively be out of service until a new namenode could be brought online!

NameNode and Secondary NameNode was a good combo

But when the NameNode failed:

HDFS became out of access

Data Nodes could not be reached

Copying metadata from Secondary was slow



HDFS High Availability

- We need a new primary namenode with one of the filesystem metadata replicas and to configure datanodes and clients to use this new namenode.
- The new namenode:
 - Loads its namespace image into memory
 - Replays its edit log
 - Receives enough block reports from the datanodes to leave safe mode.
- On large clusters (many files and blocks), the time it takes for a namenode to start from cold can be 30 minutes or more.
- The long recovery time is a problem for routine maintenance - since unexpected failure of the namenode is so rare

HDFS High Availability


- Hadoop 2: adds support for HDFS high availability (HA).
- A pair of namenodes in an active-standby configuration.
- In the event of the failure of the active namenode, the standby takes over its duties to continue servicing client requests without a significant interruption.



HDFS High Availability


- Namenodes must use **highly available shared storage** to share the edit log.
- When a standby comes up, **it reads up to the end of the shared edit log** to synchronize its state with the active namenode
- Then continues to read new entries as they are written by the active namenode.
- **Datanodes must send block reports to both namenodes** because the block mappings are stored in a namenode's memory
- Clients must be configured to handle namenode failover
- The secondary namenode's role is consumed by the standby node

Components






HDFS NameNode HA Architecture



Apache Hadoop 2.0 and YARN



Component	Description
Active NameNode	Master node that manages the file system namespace and client access. It knows where every block is stored across DataNodes.
Standby NameNode	A hot backup of the Active NameNode that keeps an up-to-date copy of metadata. It can instantly take over if the Active fails.
Secondary NameNode	Checkpointing node (not a backup) - periodically merges Active's edits and fsimage to keep metadata compact (through slow HTTP requests)
DataNodes	Worker nodes that store actual data blocks. They regularly send heartbeats and block reports
JournalNodes (for HA)	Used in HA setups to synchronize metadata in real-time between Active and Standby. They store edit logs in a quorum (usually 3 or 5 nodes – has to be odd number).

⚙ Workflow:


1. Cluster starts:
 - Both NameNodes (Active & Standby) load same `fsimage` from shared storage or prior synchronization.
2. Active NameNode:
 - Handles all client requests (reads & writes).
 - Writes every metadata change (edit) to the JournalNodes quorum.
3. Standby NameNode:
 - Monitors JournalNodes.
 - Continuously reads edits and applies them to its own namespace image.
 - Maintains a real-time, synchronized copy of Active's metadata.
4. ZooKeeper Failover Controller (ZKFC):
 - Watches both NameNodes' health.
 - If Active fails → ZKFC automatically triggers failover:
 - Demotes failed Active.
 - Promotes Standby to new Active.
5. Clients are automatically redirected to the new Active NameNode.
6. **High Availability achieved** — no downtime, no manual recovery.
7. But still only one node (Active) serves read requests → potential bottleneck in read-heavy workloads.

All good but:

Only Active was serving both read and write requests

Became a real bottleneck in big clusters

So, now, let's keep an observer node for read.
It is a read-only NameNode
Slightly lags behind Active!





⚙ Workflow:

1. Startup phase:

- Active, Standby, and Observer(s) all load the same `fsimage`.
- All connect to **JournalNodes**.

2. Active NameNode:

- Handles **all write and coordination operations**.
- Writes every metadata change (edit) to JournalNodes.

3. Standby NameNode:

- Continuously reads and applies new edits from JournalNodes.
- Ready to **take over automatically** if Active fails.

4. Observer NameNode(s):

- Also read the same edit logs from JournalNodes asynchronously.
- Keep an *almost up-to-date* copy of namespace.
- Can serve **read-only client requests** (`ls`, `stat`, `getBlockLocations`).

5. Client behavior:

- Read operations → sent to **Observer** (for better performance).
- Write operations → sent to **Active**.
- If Observer is lagging or unavailable, client falls back to Active.

6. Consistency model:

- Slight lag possible (eventual consistency for reads).
- Clients requiring "read-your-write" consistency can request Active explicitly.

7. Failover:

- ZKFC still monitors Active and Standby only.
- Observers remain read-only — they don't take over automatically.

8. Result:


- High Availability (like Hadoop 2.x)
- Plus **Read Scalability** (thanks to Observers)
- Perfect for **large-scale, read-heavy clusters** (e.g., Hive, Spark, Impala).

Note:

- You can also back up the files that make up the persistent state of namenode metadata.
- `hdfs dfsadmin -fetchImage /backup/fsimage`
- `hdfs dfsadmin -saveNamespace`
- If both Active, Standby and Observers are lost, you can restore the cluster using this offline backup.




Block Caching






Components





Block Caching

- A datanode reads blocks from disk
- For frequently accessed files the blocks may be explicitly cached in the datanode's memory, in an off-heap block cache.
- A block is cached in only one datanode's memory, although the number is configurable on a per-file basis.
- Job schedulers (for MapReduce, Spark) take advantage of cached blocks by running tasks on the datanode where a block is cached
- A small lookup table used in a join is a good candidate for caching, for example.





HDFS Federation

- HDFS federation, allows a cluster to scale by adding namenodes, each of which manages a portion of the filesystem namespace.
- For example, one namenode might manage all the files rooted under /user, say, and a second name- node might handle files under /share.
- Hadoop Federation allows **multiple independent NameNodes**, each managing a separate part of the namespace (called a “namespace volume”)
- This gives both **scalability** and **partial redundancy** — if one NameNode fails, others are unaffected




HDFS Federation

- Under federation, each namenode manages a namespace volume: **metadata + block pool** containing all the blocks for the files in the namespace.
- Namespace volumes are independent:
 - Namenodes do not communicate with one another
 - Failure of one namenode does not affect availability of namespaces managed by other namenodes.
- Block pool storage is not partitioned
 - Datanodes register with each namenode in the cluster and store blocks from multiple block pools




HDFS Federation





HDFS Federation

HDFS Federation Architecture





Uber

We use **Hadoop for both batch and streaming analytics** across a wide range of use cases, such as fraud detection, machine learning, and ETA calculation.

Uber's Data Infrastructure team massively overhauled scaling HDFS by implementing new adjustments:

- [View File System](#) (ViewFs)
- Frequent HDFS version upgrades
- NameNode [garbage collection](#) tuning
- Limiting the number of small files that filter through the system
- An HDFS load management service
- A read-only NameNode replica



» General

- [Overview](#)
- [Single Node Setup](#)
- [Cluster Setup](#)
- [Commands Reference](#)
- [FileSystem Shell](#)
- [Compatibility Specification](#)
- [Downstream Developer's Guide](#)
- [Admin Compatibility Guide](#)
- [Interface Classification](#)
- [FileSystem Specification](#)

» Common

- [CLI Mini Cluster](#)
- [Fair Call Queue](#)
- [Native Libraries](#)
- [Proxy User](#)
- [Rack Awareness](#)
- [Secure Mode](#)
- [Service Level](#)

ViewFs Guide

- [Introduction](#)
- [The Old World \(Prior to Federation\)](#)
 - [Single Namenode Clusters](#)
 - [Pathnames Usage Patterns](#)
 - [Pathname Usage Best Practices](#)
- [New World – Federation and ViewFs](#)
 - [How The Clusters Look](#)
 - [A Global Namespace Per Cluster Using ViewFs](#)
 - [Pathname Usage Patterns](#)
 - [Pathname Usage Best Practices](#)
 - [Renaming Pathnames Across Namespaces](#)
- [Multi-Filesystem I/O with Nfly Mount Points](#)
 - [Basic Configuration](#)
 - [Advanced Configuration](#)
 - [Network Topology](#)
 - [Example Nfly Configuration](#)
 - [How Nfly File Creation works](#)
 - [FAQ](#)
- [Don't want to change schema or difficult to copy mount-table configurations to all clients?](#)

Provides a way to manage multiple Hadoop file system namespaces

It is particularly useful for clusters having multiple namenodes, and hence multiple namespaces in the HDFS Federation

Uber


HDFS was designed to support thousands of nodes within a single cluster. With enough hardware, scaling to over 100 petabytes of raw storage capacity in one cluster can be easily—and quickly—achieved.

For Uber, however, the rapid growth of our business made it **difficult to scale reliably without slowing down data analysis** for our thousands of users making **millions** of Hive or Presto queries each week.






MicroStrategy



Superset



Presto Cluster



Object Storage



RDBMS Storage



NoSQL Sources



EMR






Uber

Presto accounts for more than half of the access to HDFS, and 90 percent of Presto queries take ~100 seconds to process - **If our HDFS infrastructure is overloaded, Presto queries in the queue pile up (completion delays)** - Also data should be available immediately.

We engineered **ETL to occur in the same clusters where users run queries** to reduce replication latency - but results in the generation of small files to accommodate frequent writes, which further clogs the queue.






Uber

Another one! **Multiple teams require a large percentage of stored data**, making it impossible to split clusters by use case or organization without duplications - decrease efficiency and increase costs!

The root of these slowdowns:

performance and throughput of the NameNode - the directory tree of all files in the system that tracks where data files are kept.


All metadata stored in NameNode - client requests to an HDFS cluster must first pass through it.





Uber

We started experiencing high NameNode remote procedure call (RPC) queue time - NameNode queue time could exceed 500 milliseconds per request - every single HDFS request waited for at least half a second in the queue (normal = 10 ms)





Uber


Inspired by a similar effort at Twitter, we utilized [View File System](#) (ViewFs) to split our HDFS into multiple physical namespaces and used ViewFs mount points to present a single virtual namespace to users,

Separated our HBase from HDFS cluster of the YARN and Presto operations.

Greatly reduced the load on the main cluster

Reduced the HBase cluster restart from hours to minutes.

Typical Mount Table for each Cluster




Uber

Implemented **separate HDFS clusters behind ViewFS** instead of HDFS Federation –

Now, HDFS upgrades gradually rolled out to minimize the risk of large-scale outages; complete isolation also helps improve system reliability

Downside: separate HDFS clusters = slightly higher operational costs.



Uber

A second solution for our scaling challenges was to **upgrade our HDFS to keep up with the latest release versions.**

Upgrading HDFS can be **risky** because it can cause **downtime, performance degradation, or data loss.**

To counter these possible issues, we spent several months validating 2.8.2 before deploying it in production.



Uber

Run different versions of YARN and HDFS on the same servers at the same time - critical to scaling.

Since both YARN and HDFS are part of Hadoop, they are normally upgraded together.

However, major YARN upgrades take much longer to fully validate and roll out

Because some production applications running on YARN might need to be updated due to YARN API changes or JAR dependency conflicts between YARN and these applications.



Garbage collection (GC) tuning has also played an important role in our optimization and given us much needed breathing room as we continue to scale out our storage infrastructure.

We prevent long GC pauses by forcing Concurrent Mark Sweep collectors (CMS) to do more aggressive old generation collections

For reference, our customized GC-related Java Virtual Machine (JVM) arguments for NameNode with 160GB heap size are:

- XX:+UnlockDiagnosticVMOptions
- XX:ParGCCardsPerStrideChunk=32768 -XX:+UseParNewGC
- XX:+UseConcMarkSweepGC -XX:+CMSConcurrentMTEnabled
- XX:CMSInitiatingOccupancyFraction=40
- XX:+UseCMSInitiatingOccupancyOnly
- XX:+CMSParallelRemarkEnabled -XX:+UseCondCardMark
- XX:+DisableExplicitGC




Uber

NameNode loads all file metadata in memory, **storing small files increases memory pressure on the NameNode** - increase of read RPC calls for accessing same amount of data when clients are reading files

Built new ingestion pipelines based on our Hoodie library that generates much bigger files than those created by our original data pipelines.

‘stitcher’ - merges together small files into larger ones, mostly greater than 1GB in size.






Uber

Second, we set a strict namespace quota on Hive databases and application directories - created a self-service tool for users to manage quotas within their organizations - allocated at the ratio of 256MB per file - encourage users to optimize their output file size.

Another one! Detecting which applications are causing unusually large loads, and from there, taking quick action to fix them.

We built an in-house HDFS load management service, referred to as **Spotlight**, for this purpose.

In Spotlight, an audit log is streamed from active NameNode and processed in real time through a backend based on Flink and Kafka - output is shown on dashboard - automatically disable accounts or kill tasks causing slowdown.



Switch View (Itemized / Aggregated)


time	rpcquetime	cmd	count	ugi	directory
		contentSummary	332	hive	/user/hive/.sparkStaging
02-26 18:10:42 – 02-26 18:11:13		create	24973	hive	/tmp/hive_stagingdir/hive_hive_2018-02-27_02-02-28_397_5474197884312459232-59776
		create	22438	hive	/tmp/hive_stagingdir/hive_hive_2018-02-27_02-01-30_604_5744798450550671345-59763
		create	19347	hive	/tmp/hive_stagingdir/hive_hive_2018-02-27_02-08-11_715_2470669021781026765-59098
		create	8172	hive	/tmp/hive_stagingdir/hive_hive_2018-02-27_02-01-52_950_8449284600356918220-58896
		create	3845	hive	/tmp/hive_stagingdir/hive_hive_2018-02-27_02-01-13_545_7766103451547089863-22982
		create	1764	hive	/tmp/hive_stagingdir/hive_hive_2018-02-27_02-05-02_486_3190944843817077058-6249
		create	1103	hive	/tmp/hive_stagingdir/hive_hive_2018-02-27_02-07-55_055_25282529972543290-22501




Uber

Development of **Observer NameNode** -
read-only NameNode replica - reduce
the load on the active NameNode
cluster.

Because **more than half of HDFS RPC volume comes from read-only Presto queries**, Observer NameNodes help us scale the overall NameNode throughput by close to 100 percent in the first release.






Uber

Layer your solutions: Scalability improvements like Observer NameNode or splitting HDFS into more clusters takes significant effort - Short-term measures like GC tuning and merging smaller files into larger ones via our stitcher gave us a lot of breathing room to develop long-term solutions.

Bigger file sizes are better: Since small files are a threat to HDFS, it is better to tackle them earlier rather than later.

Participate in the community: Hadoop has been around for more than 10 years and its community is more active than ever before, leading to scalability and functionality improvements introduced in nearly every release.





ViewFs - scale out HDFS when subclusters become overloaded – But then **client configuration changes are required every time one adds or replaces a new mount point** - affects production workflows.

Microsoft's Router-based Federation - HDFS 2.9 release - natural extension to a ViewFs-based partitioned federation.

Adds a layer of software capable of centralizing HDFS namespaces –

By offering the same interface (a combination of RPC and WebHDFS), **gives users transparent access to any subcluster** - subclusters manage their data independently.



Uber

Importance of reducing storage costs - users access recent data (“hot” data) more frequently than old data (“warm” data).

Moving old data to a separate, less resource-intensive tier will significantly reduce our storage costs.

HDFS Erasure Coding, Router-based Federation, high density (more than 250TB) hardware, and a data mover service (to handle moving data between “hot” tier clusters and “warm” tier clusters) are key components of our forthcoming tiered storage design.



VectorStock

VectorStock.com/6366282