

# Big Data Analytics

Fall 2025

## Lecture 10

Dr. Tariq Mahmood



# Data Analysts with Hadoop



# Challenges of MapReduce



- Requires complex, verbose Java (python using PySpark these days)
- Complex: Understand distributed computing
- Imperative programming: define step-by-step procedures
- No support for SQL (motivation for hive)
- Optimized for batch processing – latency issues
- Job execution stack too long for simple queries (remember uber?)
- No schema management
- No connectivity with BI tools (can connect through Hive, Presto)
- Maintenance burden



**... Enter Hive!**



# The Base of Hive

- Data warehouse infra on top of Hadoop
- Enables data summarization
- Enables ad-hoc queries - select \* type queries :)
- Initially developed by Facebook.
- Hive stores data in HDFS
- Supports SQL like Query Language : HiveQL
- HQL statements are broken down by the Hive service into MapReduce jobs
  - For MR execution engine only
  - Tez engine (DAG-based), Spark engine (Spark-based) and Long Live and Process (LLAP) (in-memory caching)





<b>Engine</b>	<b>Description</b>	<b>Enabled by</b>
<b>MapReduce</b>	Original, stable, slow	Default in older Hadoop
<b>Apache Tez</b>	DAG-based engine — faster, optimized	hive.execution.engine=tez
<b>Apache Spark</b>	Uses Spark as execution engine	hive.execution.engine=spark
<b>LLAP (Live Long and Process)</b>	In-memory caching for low-latency	Part of Hive-on-Tez

<b>Step</b>	<b>Description</b>
<b>1. Parse</b>	Hive parses HQL into an Abstract Syntax Tree (AST).
<b>2. Semantic Analysis</b>	Validates schema & columns via Metastore.
<b>3. Logical Plan</b>	Builds a logical operator tree (select, filter, join, etc.).
<b>4. Optimization</b>	Applies rule-based or cost-based optimizations.
<b>5. Physical Plan</b>	Converts logical plan to physical plan → MapReduce DAG / Tez DAG / Spark DAG.
<b>6. Execution</b>	Submitted to YARN for distributed execution.

# Organization

- Data in Hive is organized into Tables
- Work with data inside HDFS
- Tables:
  - Data: Files in HDFS
  - Schema: In the form of metadata stored in RDBMS (logical layer)
  - Have a corresponding HDFS directory
  - Data in a table is serialized
- Supports primitive types and also nestable collections: Array and Map(Key Value pair)



# Data Model

- Tables
  - Each table has a corresponding HDFS directory
  - Hive provides built-in serialization formats (lazy serialization – delayed until needed)
- **Partitions:** Each table can have 1 or more horizontal partitions
  - Example: Table T in the directory : /wh/T.
  - If T is partitioned on columns ds = ‘20090101’, and ctry = ‘US’, will be stored /wh/T/ds=20090101/ctry=US.
- **Buckets:** Data in each partition may be divided into buckets based on the hash of a column in the table
  - Each bucket is stored as a file in the partition directory

```
CREATE TABLE ecommerce_sales (
    order_id STRING,
    category STRING,
    product_id STRING,
    price FLOAT,
    payment_type STRING
)
PARTITIONED BY (country STRING, order_date DATE)
STORED AS PARQUET;
```

```
/warehouse/ecommerce_sales/country=Pakistan/order_date=2025-11-01/
/warehouse/ecommerce_sales/country=Pakistan/order_date=2025-11-02/
/warehouse/ecommerce_sales/country=USA/order_date=2025-11-02/
```

```
SELECT SUM(price)
FROM ecommerce_sales
WHERE country = 'Pakistan' AND order_date = '2025-11-02';
```

Only 1 partition will be read – no need to scan the whole dataset

```
CREATE TABLE ecommerce_sales_bucketed (
    order_id STRING,
    customer_id STRING,
    category STRING,
    product_id STRING,
    price FLOAT,
    payment_type STRING
)
PARTITIONED BY (country STRING)
CLUSTERED BY (customer_id) INTO 8 BUCKETS
STORED AS ORC;
```

```
/warehouse/ecommerce_sales_bucketed/country=Pakistan/000000_0
/warehouse/ecommerce_sales_bucketed/country=Pakistan/000001_0
...
/warehouse/ecommerce_sales_bucketed/country=Pakistan/000007_0
```

```
SELECT customer_id, COUNT(order_id)
FROM ecommerce_sales_bucketed
WHERE country = 'Pakistan'
GROUP BY customer_id;
```

All records for the same cst  
always go to the same bucket  
file ( $=\text{hash}(\text{cid}) \bmod 8$ )  
**A single cst is completely in 1  
bucket – not spread out.**

Hive assigns **each bucket to  
one reducer only**

**Each bucket file corresponds to  
a  $\text{hash}(\text{cid}) \bmod 8$**


**Data pre-hashed by cid, so Hive  
parallelizes aggregations and joins  
using bucket files.**




```
CREATE TABLE page_view(viewTime INT, userid BIGINT,
    page_url STRING, referrer_url STRING,
    friends ARRAY<BIGINT>, properties MAP<STRING, STRING>
    ip STRING COMMENT 'IP Address of the User')
    COMMENT 'This is the page view table'
PARTITIONED BY(dt STRING, country STRING)
CLUSTERED BY(userid) SORTED BY(viewTime) INTO 32 BUCKETS
ROW FORMAT DELIMITED
    FIELDS TERMINATED BY '1'
    COLLECTION ITEMS TERMINATED BY '2'
        MAP KEYS TERMINATED BY '3'
STORED AS SEQUENCEFILE;
```

# Hive Key Principles


Enables users to query data as if they were using a relational database - much simpler than writing raw MapReduce code



# Hive Key Principles




# Hive Key Principles



External/Internal Tables, Schema-on-read, Join/Group-by, UDF, Partitioning, Bucketing, Direct load, Spark integration,

# Hive Key Principles



**TRANSFORM:** incorporate custom **mapper** and **reducer** scripts directly into Hive queries.

# Hive is Not...



- Hive designed for batch processing of large datasets
- Not an OLTP or real-time system
- **Latency** (Time to start & finish one query) is high compared to RDBMS
- **Throughput** (Data processed per batch) is high compared to RDBMS
- Even when dealing with relatively small data ( <100 MB )

I Am, I Am Not



# Nuts and Bolts


- Supports primitive types, and Associative Arrays (maps, dictionaries, hash tables) Lists, Struct.





# Nuts and Bolts

- HQL supports DDL
  - Define or modify the structure - Create, alter, drop, truncate
  - Automatically commits
- HQL supports DML
  - Don't change the database schema but rather modify the actual data within tables.
  - Insert, update, delete, select
- HQL has limited equality and join predicates (inner, left and right outer, full) - and has no inserts on existing tables.
- But improving in newer versions





Requirement	Description
Format	ORC or Parquet recommended
Bucketing	Required for full ACID tables
Property	'transactional'='true'
Execution engine	Tez or Spark (not MR)
HiveServer2 + Metastore	Must be configured with concurrency control

:)

Operation	Supported?	Notes
INSERT INTO	Yes	Appends new files to table (always supported)
INSERT OVERWRITE	Yes	Replaces table or partition data
UPDATE / DELETE	Yes (only in ACID tables)	Requires transactional setup
MERGE	Yes (ACID only)	For upsert operations



# Serialize - Deserialize

- Convert data into a format suitable for storage or transmission - and reconstruct back to original form
- Purpose of Ser: **persist** the data's state
- Purpose of DeSer: **reconstruct** original state for application usage
- Allows complex data structures (lists, dictionaries) to be saved as a single entity
- Considerations:
  - Overhead
  - Data Integrity (for non-ACID tables)

```
{  
  "id": 1,  
  "name": "John Doe",  
  "email": "john.doe@example.com"  
}
```



# Serialize - Deserialize

```
class Person:  
    def __init__(self, name, age):  
        self.name = name  
        self.age = age  
  
    def __repr__(self):  
        return f"Person(name={self.name}, age={self.age})"  
  
# Create an instance of Person  
person = Person("John Doe", 30)  
print("Original object:", person)
```

```
# Serialize the object  
serialized_person = pickle.dumps(person)  
  
# Show the serialized binary data  
print("Serialized object (binary):", serialized_person)  
  
# Deserialize the object  
deserialized_person = pickle.loads(serialized_person)  
  
# Show the deserialized object  
print("Deserialized object:", deserialized_person)
```


Original object: Person(name=John Doe, age=30)

Serialized object (binary): b'\x80\x04\x95...'

Deserialized object: Person(name=John Doe, age=30)



# Hive Architecture




**HADOOP**  
(MAP-REDUCE + HDFS)

Job  
Tracker      Name  
Node


Data Node  
+  
Task Tracker

# Hive Architecture




- External Interfaces:
  - Web UI : Management
  - Hive CLI : Run Queries, Browse Tables etc
  - API : JDBC, ODBC
- Metastore: System catalog - contains metadata about Hive tables
- **Driver:** Manages life cycle of a HiveQL statement during compilation, optimization and execution

# Hive Architecture



- Compiler : translates HiveQL statement into a plan which consists of a DAG of map-reduce jobs – Why DAG?
- Database: is a namespace for tables (meaning?)
- **Table**: Metadata for table contains list of columns and their types, owner, storage and SerDe information - Also contains any user supplied key and value data.
- **Partition**: Each partition can have its own columns and SerDe and storage information.

# Why DAG for queries?



Dependency Management

Optimize – eliminate unnecessary steps

Parallel execution

Fault tolerance – run only the dependent nodes when a task fails

Effective resource allocation – as we know dependencies

### Multi CTEs - DAG in SQL

Last edit was made 0 seconds ago · by [REDACTED]

```
1 WITH
2   cte1 AS (
3     SELECT
4       a,
5       b
6     FROM mytable
7   ),
8
9   cte2 AS (
10    SELECT c
11    FROM table2
12  ),
13
14   cte3 AS (
15     SELECT
16       ct1.a,
17       cte2.c
18     FROM cte1
19     INNER JOIN cte2 ON cte1.b = cte2.c
20  ),
21
22 final_result AS (
23   SELECT
24     cte1.a,
25     cte3.c
26   FROM cte1
27   INNER JOIN cte3 ON cte1.a = cte3.a
28 )
29
30 SELECT *
31 FROM final_result;
32
```

mysql


Properties

Open in a new page



Share

...



Doc

Comment


History

Snippet

Linter

Lineage


# Bottom




Top

Figure 2: Query plan with 3 map-reduce jobs for multi-table insert query

## *System for querying and managing structured data*





**External interface:**  
Both user interface like command line (cli) and web UI

Thrift is a framework for cross-language services, where a server written in one language (like Java) can also support clients in other languages.


Metastore is the system catalog. All other components of Hive interact with metastore

The Driver manages the life cycle (statistics) of a HiveQL statement during compilation, optimization and execution



# CLI

- There are several ways to interact with Hive, including some popular graphical user interface but CLI is sometimes preferable (**Beeline**)
- CLI allows creating, inspecting schema and query tables, etc.
- All commands and queries go to the Driver, which complies, optimizes and executes queries usually with MapReduce jobs.
- Hive communicates with Job Tracker to initiate the MapReduce job.






# Metastore

- The system catalog which contains **metadata** about the tables stored in Hive
- This data is specified during table creation and reused every time the table is referenced in HiveQL
- Contains the following objects:
  - Database : the namespace for tables
  - Table: metadata for table contains list of columns and their types, owners, storage and SerDe information
  - Partition: each partition can have its own columns and SerDe and storage information




# Hive Metastore (Cloudera)



Thrift: Applications in Different programming languages can communicate, Ser and DeSer, Scalable

Hive Metastore- centralizes the metadata



Centralized security and policy management

Access control – data auditing – data masking

# Compile

- Converts DDL-DML to a plan.
- Parser: transforms a query to a parse tree representation
- Semantic analyzer: parse tree to internal (block-based) representation
- Plan generator: internal representation to a logical plan
- Optimizer: multiple passes over plan and rewrites
  - Combine multiple joins sharing join key into single multiway join - single MR!
  - Add **repartition** operators (redistribute data to balance workload)
  - **Prune columns early** (select req. columns as early as possible)
  - **Push query predicates closer** to the table scan operators (predicate pushdown – why?)

It's time to build.

# Compile

It's time to **build.**

- Physical Plan generator: converts logical plan into physical plan = DAG of MR jobs



# HiveQL

- SQL like language: HiveQL
- DDL : to create tables with specific serialization formats
- DML : load and insert to load data from external sources and insert query results into Hive tables
- Do not support updating and deleting rows in existing tables
- Supports Multi-Table insert
- Supports Select, Project, Join, Aggregate
- Supports Union all and Sub-queries in the From clause



# HiveQL

- Can be extended with custom functions (UDFs)
- User Defined Transformation Function(UDTF)  
<https://cwiki.apache.org/confluence/display/Hive/DeveloperGuide+UDTF>
- User Defined Aggregation Function (UDAF)
- Users can embed custom map-reduce scripts written in any language using a simple row-based streaming interface

- Running time example: Status Meme

When Facebook users update their status, the updates are logged into flat files in an NFS directory /logs/status\_updates

status\_updates(userid int, status string, ds string)  
using a load statement like below.

```
LOAD DATA LOCAL INPATH '/logs/status_updates'  
INTO TABLE status_updates PARTITION (ds='2009-03-20')
```

```
FROM (SELECT a.status, b.school, b.gender  
      FROM status_updates a JOIN profiles b  
        ON (a.userid = b.userid and  
            a.ds='2009-03-20' )  
    ) subq1
```

```
INSERT OVERWRITE TABLE gender_summary  
          PARTITION(ds='2009-03-20')
```

```
SELECT subq1.gender, COUNT(1) GROUP BY subq1.gender
```

```
INSERT OVERWRITE TABLE school_summary  
          PARTITION(ds='2009-03-20')
```

```
SELECT subq1.school, COUNT(1) GROUP BY subq1.school
```

Compute daily statistics on the frequency of status updates based on gender and school



# Advantages – Child's play

- Fast response
- Scalable and extensible
- Thrift
- Ser/De
- Flexible schema



# Advantages – Child's play

- No row-level insert, update or delete operations – load it!
- Four file formats: TEXTFILE, SEQUENCEFILE, ORC and RCFILE.
- Example: 'NASDAQ\_daily\_prices\_B.csv' a log file of stocks record of NASDAQ.
- exchange,stock\_symbol,date,stock\_price\_open,stock\_price\_high,stock\_price\_low,stock\_price\_close,stock\_volume,stock\_price\_adj\_close
- NASDAQ,BBND,2010-02-08,2.92,2.98,2.86,2.96,483800,2.96  
NASDAQ,BBND,2010-02-05,2.85,2.94,2.79,2.93,884000,2.93  
NASDAQ,BBND,2010-02-04,2.83,2.88,2.78,2.83,1333300,2.83



```
hive> CREATE TABLE IF NOT EXISTS stocks (
    exchange STRING,
    symbol STRING,
    ymd STRING,
    price_open FLOAT,
    price_high FLOAT,
    price_low FLOAT,
    price_close FLOAT,
    volume INT,
    price_adj_close FLOAT)
ROW FORMAT DELIMITED FIELDS
TERMINATED BY ',';
```

- Create a database:

```
hive> CREATE DATABASE financials;
```

or

```
hive> CREATE DATABASE IF NOT EXISTS financials;
```

- Describe table:

```
hive> DESCRIBE DATABASE financials;
```

OK

Financials

```
hdfs://localhost:54310/user/hive/warehouse/financials.db
```

- Use database:

```
hive> USE financials;
```

- Drop database:

```
hive> DROP DATABASE IF EXISTS financials;
```



# Some Facts

- Use LOAD DATA to import data into a Hive table
- Hive>Load Data LOCAL INPATH '/home/sunny/EmployeeDetails.txt' INTO TABLE Employee
- Use the word OVERWRITE to write over a file of the same name
- We can Load data from Local file system by using LOCAL keyword
- Inserting Data into new table by using SELECT statement
- For Example, INSERT OVERWRITE SELECT \* FROM Employee



Operation	Command Syntax
See current tables	Hive>Show TABLES
Check the table name	Hive>Describe <Table_Name>
Change the table name	Hive>Alter Table <table_Name> Rename to mytab
Add a column	Hive> Alter Table <table_Name> ADD COLUMNS (MyID String)
Drop a partition	Hive>Alter Table <table_Name> DROP PARTITION (Age>70)



# Some Facts

- WHERE Clause
- UNION All and DISTINCT
- GROUP BY and HAVING
- LIMIT Clause
- Hive Supports Sub-Queries but only in FROM Clause
- JOINS , ORDER BY, SORT BY



# Output Data

- Output data produced by Hive is structured, typically stored in a relational database.
- For cluster, MySQL or similar relational database is required.
- The result tables then can be manipulated using HiveQL in the similar way of SQL to relational database.



```
hive> LOAD DATA LOCAL INPATH  
'/Users/nqt289/Desktop/NASDAQ_daily_prices_B.csv'  
    > OVERWRITE INTO TABLE stocks;
```

Copying data from  
file:/Users/nqt289/Desktop/NASDAQ\_daily\_prices\_B.csv

Copying file:  
file:/Users/nqt289/Desktop/NASDAQ\_daily\_prices\_B.csv

Loading data to table mydb.stocks

Deleted

hdfs://localhost:54310/Users/nqt289/Desktop/NASDAQ\_  
daily\_prices\_B.csv

OK

Time taken: 0.231 seconds



```
hive> SELECT * FROM STOCKS WHERE price_open='2.92';
```

Total MapReduce jobs = 1

Launching Job 1 out of 1

Number of reduce tasks is set to 0 since there's no reduce operator

Starting Job = job\_201403311509\_0003, Tracking URL = [http://localhost:50030/jobdetails.jsp?jobid=job\\_201403311509\\_0003](http://localhost:50030/jobdetails.jsp?jobid=job_201403311509_0003)

```
Kill Command = /Users/nqt289/hadoop-0.20.2/bin/..bin/hadoop job -Dmapred.job.tracker=localhost:54311 -kill job_201403311509_0003
```

## Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0

2014-03-31 15:39:20,577 Stage-1 map = 0%, reduce = 0%

2014-03-31 15:39:23,597 Stage-1 map = 100%, reduce = 0%

2014-03-31 15:39:26,625 Stage-1 map = 100%, reduce = 100%

Ended Job = job\_201403311509\_0003

## MapReduce Jobs Launched:

Job 0: Map: 1 HDFS Read: 21998523 HDFS Write: 5166 SUCCESS

Total MapReduce CPU Time Spent: 0 msec

OK

NASDAQ BBND 2010-02-08 2.92 2.98 2.86 2.96 483800  
2.96

NASDAQ BTFG 2009-12-21 2.92 2.92 2.75 2.79 15100

NASDAQ BJCT 2004-04-21 2.92 2.98 2.9 2.98 3200

NASDAQ BJCT 2004-04-20 2.92 3.0 2.92 2.95 27900

# Primitive Data Types

Type	Comments
TINYINT, SMALLINT, INT, BIGINT	1, 2, 4 and 8-byte integers
BOOLEAN	TRUE/FALSE
FLOAT, DOUBLE	Single and double precision real numbers
STRING	Character string
TIMESTAMP	Unix-epoch offset <i>or</i> datetime string
DECIMAL	Arbitrary-precision decimal
BINARY	Opaque; ignore these bytes

# Complex Data Types

Type	Comments
STRUCT	A collection of elements If S is of type STRUCT {a INT, b INT}: S.a returns element a
MAP	Key-value tuple If M is a map from 'group' to GID: M['group'] returns value of GID
ARRAY	Indexed list If A is an array of elements ['a','b','c']: A[0] returns 'a'

# HiveQL Limitations

- Missing large parts of full SQL specification:
  - HAVING clause in SELECT (must be used after Group By – !directly)
  - Correlated sub-queries (use them with JOIN)
  - Updatable or materialized views (updatable views not possible)
  - Stored procedures (not allowed; rather use UDFs)

```
SELECT department, COUNT(employee_id) AS  
employee_count  
FROM employees  
GROUP BY department  
HAVING employee_count > 5;
```





```
SELECT employee_id, salary
FROM employees e
WHERE salary > (
    SELECT AVG(salary)
    FROM employees
    WHERE department_id = e.department_id
);
```



```
SELECT e.employee_id, e.salary
FROM employees e
JOIN (
    SELECT department_id, AVG(salary) AS avg_salary
    FROM employees
    GROUP BY department_id
) dept_avg
ON e.department_id = dept_avg.department_id
WHERE e.salary > dept_avg.avg_salary;
```



```
WITH dept_avg AS (
    SELECT department_id, AVG(salary) AS avg_salary
    FROM employees
    GROUP BY department_id
)
SELECT e.employee_id, e.salary
FROM employees e
JOIN dept_avg
ON e.department_id = dept_avg.department_id
WHERE e.salary > dept_avg.avg_salary;
```



```
CREATE MATERIALIZED VIEW sales_summary  
AS  
SELECT product_id, SUM(sales_amount) AS total_sales  
FROM sales  
GROUP BY product_id;
```

```
ALTER MATERIALIZED VIEW sales_summary REBUILD;
```

# Hive Metastore



- Stores Hive metadata
- Default metastore database uses Apache Derby
- Default: **embedded, single-user database** in local or testing environments
- **Only one active connection** at a time
- For production, replace with Postgres, MySQL (configure `hive-site.xml`)

# Hive Metastore

- **Embedded:** Metastore runs within the Hive service process – uses embedded Derby - for unit tests, development, single-user
- **Local:** Each Hive service has its own metastore process - client connects to the metastore directly – shared external db – for multi-user environment
- **Remote:** Metastore runs as a separate standalone service - Each Hive client connects to the metastore server, which connects to the metadata database itself.



## Embedded Metastore



# Hive Warehouse

- Hive tables are stored in the Hive “warehouse”
  - Default HDFS location: /user/hive/warehouse
- Tables are stored as sub-directories in the warehouse directory
- Partitions are subdirectories of tables
- External tables are supported in Hive
- The actual data is stored in flat files

## File Size Comparison Across Encoding Methods

Dataset: TPC-DS Scale 500 Dataset

**585 GB**  
(Original Size)

Encoded with  
**Text**

**505 GB**  
(14% Smaller)

Encoded with  
**RCFile**

**Impala**  
**221 GB**  
(62% Smaller)

Encoded with  
**Parquet**

**Hive 12**  
**131 GB**  
(78% Smaller)

Encoded with  
**ORCFile**

- Larger Block Sizes
- Columnar format arranges columns adjacent within the file for compression & fast access

# Hive Schemas

- Hive is **schema-on-read**
  - Schema is only enforced when the data is read (at query time)
  - Allows greater flexibility: same data can be read using multiple schemas
- Contrast with an RDBMS, which is schema-on-write
  - Schema is enforced when the data is loaded
  - Speeds up queries at the expense of load times

# Create Table Syntax

```
CREATE TABLE table_name
  (col1 data_type,
   col2 data_type,
   col3 data_type,
   col4 datatype )
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
STORED AS format_type;
```

# Simple Table

```
CREATE TABLE page_view
(viewTime INT,
 userid BIGINT,
 page_url STRING,
 referrer_url STRING,
 ip STRING COMMENT 'IP Address of the User' )
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE;
```

# More Complex Table

```
CREATE TABLE employees  (
    name STRING,
    salary FLOAT,
    subordinates ARRAY<STRING>,
    deductions MAP<STRING, FLOAT>,
    address STRUCT<street:STRING,
                city:STRING,
                state:STRING,
                zip:INT>)

ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE;
```

# More About Tables

- CREATE TABLE
  - LOAD: file moved into Hive's data warehouse directory
  - DROP: both metadata and data deleted
- CREATE EXTERNAL TABLE
  - LOAD: no files moved
  - DROP: only metadata deleted
  - Use this when sharing with other Hadoop applications, or when you want to use multiple schemas on the same data

# Managed vs External

- Managed: Hive owns and manages both the metadata and the actual data stored on HDFS.
- When Hive is the main tool for managing data and when you want Hive to control the lifecycle of both data and metadata

```
CREATE TABLE employees (id INT, name STRING)  
STORED AS ORC;
```

/user/hive/warehouse/employees

# External

- Hive manages only metadata while **actual data files are outside Hive's control.**
- The data can reside in HDFS
- Hive only maintains a reference to this data
- Dropping the table doesn't delete the data

```
CREATE EXTERNAL TABLE employees (id INT, name STRING)  
STORED AS ORC  
LOCATION '/user/data/employees';
```

# External Table

```
CREATE EXTERNAL TABLE page_view_stg
(viewTime INT,
 userid BIGINT,
 page_url STRING,
 referrer_url STRING,
 ip STRING COMMENT 'IP Address of the User')
ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t'
STORED AS TEXTFILE
LOCATION '/user/staging/page_view';
```

# Partitioning

- Can make some queries faster
- Divide data based on partition column
- Use PARTITION BY clause when creating table
- Use PARTITION clause when loading data
- SHOW PARTITIONS will show a table's partitions

```
CREATE TABLE sales (
    id INT,
    product STRING,
    amount DOUBLE
)
PARTITIONED BY (year INT, month INT)
STORED AS PARQUET;
```

```
INSERT INTO TABLE sales PARTITION (year=2024, month=10)
VALUES (1, 'ProductA', 250.0),
       (2, 'ProductB', 300.0);
```

```
SELECT * FROM sales
WHERE year = 2024 AND month = 10;
```

# Bucketing

- Can speed up queries that involve sampling the data
  - Sampling works without bucketing, but Hive has to scan the entire dataset
- Use CLUSTERED BY when creating table
  - For sorted buckets, add SORTED BY
- To query a sample of your data, use TABLESAMPLE

```
CREATE TABLE employees (
    id INT,
    name STRING,
    department STRING
)
CLUSTERED BY (id) INTO 4 BUCKETS
STORED AS TEXTFILE;
```

```
INSERT INTO TABLE employees
VALUES (1, 'Alice', 'HR'),
       (2, 'Bob', 'Finance'),
       (3, 'Charlie', 'Engineering');
```

```
SELECT * FROM employees
WHERE department = 'Engineering';
```

# Browsing Tables And Partitions

Command	Comments
SHOW TABLES;	Show all the tables in the database
SHOW TABLES 'page.*';	Show tables matching the specification ( uses regex syntax )
SHOW PARTITIONS page_view;	Show the partitions of the page_view table
DESCRIBE page_view;	List columns of the table
DESCRIBE EXTENDED page_view;	More information on columns (useful only for debugging )
DESCRIBE page_view PARTITION (ds='2008-10-31');	List information about a partition

# Loading Data

- Use LOAD DATA to load data from a file or directory
  - Will read from HDFS unless LOCAL keyword is specified
  - Will append data unless OVERWRITE specified
  - PARTITION required if destination table is partitioned

```
LOAD DATA LOCAL INPATH '/tmp/pv_2008-06-8_us.txt'  
OVERWRITE INTO TABLE page_view  
PARTITION (date='2008-06-08', country='US')
```

# Inserting Data

- Use INSERT to load data from a Hive query
  - Will append data unless OVERWRITE specified
  - PARTITION required if destination table is partitioned

```
FROM page_view_stg pvs
  INSERT OVERWRITE TABLE page_view
PARTITION (dt='2008-06-08', country='US')
SELECT pvs.viewTime, pvs.userid, pvs.page_url,
       pvs.referrer_url
 WHERE pvs.country = 'US';
```

# Inserting Data

- Normally only one partition can be inserted into with a single `INSERT`
- A multi-insert lets you insert into multiple partitions

```
FROM page_view_stg pvs
INSERT OVERWRITE TABLE page_view
PARTITION ( dt='2008-06-08', country='US' )
SELECT pvs.viewTime, pvs.userid, pvs.page_url, pvs.referrer_url WHERE pvs.country = 'US'
INSERT OVERWRITE TABLE page_view
PARTITION ( dt='2008-06-08', country='CA' )
SELECT pvs.viewTime, pvs.userid, pvs.page_url, pvs.referrer_url WHERE pvs.country = 'CA'
INSERT OVERWRITE TABLE page_view
PARTITION ( dt='2008-06-08', country='UK' )
SELECT pvs.viewTime, pvs.userid, pvs.page_url, pvs.referrer_url WHERE pvs.country = 'UK';
```

# Inserting Data During Table Creation

- Use AS SELECT in the CREATE TABLE statement to populate a table as it is created

```
CREATE TABLE page_view AS  
    SELECT pvs.viewTime, pvs.userid, pvs.page_url, pvs.referrer_url  
    FROM page_view_stg pvs  
    WHERE pvs.country = 'US';
```

# Loading And Inserting Data: Summary

Use this	For this purpose
LOAD	Load data from a file or directory
INSERT	Load data from a query <ul style="list-style-type: none"><li>• One partition at a time</li><li>• Use multiple INSERTs to insert into multiple partitions in the one query</li></ul>
CREATE TABLE AS (CTAS)	Insert data while creating a table
Add/modify external file	Load new data into external table

# Sample Select Clauses

- Select from a single table

```
SELECT *
  FROM sales
 WHERE amount > 10 AND
       region = "US";
```

- Select from a partitioned table

```
SELECT page_views.*
  FROM page_views
 WHERE page_views.date >= '2008-03-01' AND
       page_views.date <= '2008-03-31'
```

# Relational Operators

- ALL and DISTINCT
  - Specify whether duplicate rows should be returned
  - ALL is the default (all matching rows are returned)
  - DISTINCT removes duplicate rows from the result set
- WHERE
  - Filters by expression
  - Does not support IN, EXISTS or sub-queries in the WHERE clause
- LIMIT
  - Indicates the number of rows to be returned

# Relational Operators

- GROUP BY
  - Group data by column values
  - Select statement can only include columns included in the GROUP BY clause
- ORDER BY / SORT BY
  - ORDER BY performs total ordering
    - Slow, poor performance
  - SORT BY performs partial ordering
    - Sorts output from each reducer

# Advanced Hive Operations

- JOIN
  - If only one column in each table is used in the join, then only one MapReduce job will run
    - This results in 1 MapReduce job:

```
SELECT * FROM a JOIN b ON a.key = b.key JOIN c ON b.key = c.key
```
    - This results in 2 MapReduce jobs:

```
SELECT * FROM a JOIN b ON a.key = b.key JOIN c ON b.key2 = c.key
```
  - If multiple tables are joined, put the biggest table last and the reducer will stream the last table, buffer the others
  - Use left semi-joins to take the place of IN/EXISTS

```
SELECT a.key, a.val FROM a LEFT SEMI JOIN b on a.key = b.key;
```

# Advanced Hive Operations

- JOIN
  - Do not specify join conditions in the WHERE clause
    - Hive does not know how to optimise such queries
    - Will compute a full Cartesian product before filtering it
- Join Example

```
SELECT
    a.ymd, a.price_close, b.price_close
FROM stocks a
JOIN stocks b ON a.ymd = b.ymd
WHERE a.symbol = 'AAPL' AND
      b.symbol = 'IBM' AND
      a.ymd > '2010-01-01';
```

# Hive Stinger

- MPP-style execution of Hive queries
- Available since Hive 0.13
- No MapReduce

# References

- <http://hive.apache.org>

# HIVE CHEAT SHEET

## Hive Basics

### Apache Hive

It is a data warehouse infrastructure based on Hadoop framework which is perfectly suitable for data summarization, analysis and querying. It uses an SQL like language called HQL (Hive query Language)

**HQL:** It is a query language used to write the custom map reduce framework in Hive to perform more sophisticated analysis of the data

**Table:** Table in hive is a table which contains logically stored data

**Hive Interfaces:**

- Hive Interfaces includes WEB UI
- Hive command line
- HD Insight (windows server)

### Components of Hive

**Meta store:** Meta store is where the schemas of the Hive tables are stored, it stores the information about the tables and partitions that are in the warehouse.

**SerDe:** Serializer, Deserializer which gives instructions to hive on how to process records

### Thrift

A thrift service is used to provide remote access from other processors

### Meta Store

This is a service which stores the metadata information such as table schemas

### Indexes

Indexes are created to the speedy access to columns in the database

Syntax: Create index <INDEX\_NAME> on table <TABLE\_NAME>

### Hive Function Meta Commands

Show functions: Lists Hive functions and operators

Describe function [function name]: Displays short description of the particular function

Describe function extended [function name]: Displays extended description of the particular function

### Hive Functions

- UDF(User defined Functions):** It is a function that fetches one or more columns from a row as arguments and returns a single value
- UDTF( User defined Tabular Functions):** This function is used to produce multiple columns or rows of output by taking zero or more inputs
- Macros:** It is a function that uses other Hive functions
- User defined aggregate functions:** A user defined function that takes multiple rows or columns and returns the aggregation of the data
- User defined table generating functions:** A function which takes a column from single record and splitting it into multiple rows

### Hive SELECT Command

```
SELECT [ALL | DISTINCT] select_expr, select_expr, ...
FROM table_reference
[WHERE where_condition]
[GROUP BY col_list]
[HAVING having_condition]
[CLUSTER BY col_list] [[DISTRIBUTE BY col_list] [SORT BY col_list]]
[LIMIT number]
;
• Select: Select is a projection operator in HiveQL, which scans the table specified by the FROM clause
• Where: Where is a condition which specifies what to filter
• Group by: It uses the list of columns, which specifies how to aggregate the records
• Cluster by, Distribute by, Sort by: Specifies the algorithm to sort, distribute and create cluster, and the order for sorting
• Limit: This specifies how many records to be retrieved
```

### Hive Data Types

#### Integral data types:

- Tinyint
- Smallint
- Int
- Bigint

#### String types:

- VARCHAR-Length(1 to 65535)
- CHAR-Length(255)

**Union type:** It is a collection of heterogeneous data types.

- Syntax: UNIONTYPE<int, double, array<string>, struct<a:int,b:string>>

#### Timestamp:

- It supports the traditional Unix timestamp with optional nanosecond precision
- Dates
  - Decimals

#### Complex types:

- Arrays: Syntax-ARRAY<data\_type>
- Maps: Syntax- MAP<primitive\_type, data\_type>
- Structs: STRUCT<col\_name : data\_type [COMMENT col\_comment], ...>

### Bucketing

It is a technique to decompose the datasets into more manageable parts

### Partitioner

Partitioner controls the partitioning of keys of the intermediate map outputs, typically by a hash function which is same as the number of reduce tasks for a job

- Partitioning:** It is used for distributing load horizontally. It is a way of dividing the tables into related parts based on values such as date, city, departments etc.

### Catalog

It is a metadata and table management system for Hadoop platform which enables storage of data in any format.


### Hive commands in HQL

**Data Definition Language(DDL):** It is used to build or modify tables and objects stored in a database. Some of the DDL commands are as follows:

- To create database in Hive: `create database <data base name>`
- To list out the databases created in a Hive warehouse: `show databases`
- To use the database created: `USE <data base name>`
- To describe the associated database in metadata: `describe <data base name>`
- To alter the database created: `alter <data base name>`

**Data Manipulation Language(DML):** These statements are used to retrieve, store, modify, delete, Insert and update data in a database

- Inserting data in a database: The Load function is used to move the data into a particular Hive table.
- `LOAD data <LOCAL> inpath <file path> into table [tablename]`
- Drop table: The drop table statements deletes the data and metadata from the table: `drop table <table name>`
- Aggregation: It is used to count different categories from the table : `Select count (DISTINCT category) from tablename;`
- Grouping: Group command is used to group the result set, where the result of one table is stored in the other: `Select <category>, sum(<amount>) from <txt records> group by <category>`
- To exit from the Hive shell: Use the command `quit`



### Operations - Performed on Hive

Function	HQL Query
To retrieve information	<code>SELECT * from _columns FROM table WHERE conditions;</code>
To select all values	<code>SELECT * FROM table;</code>
To select a particular category values	<code>SELECT * FROM table WHERE rec_name = "value";</code>
To select for multiple criteria	<code>SELECT * FROM TABLE WHERE rec1 = "value1" AND rec2 = "value2";</code>
For selecting specific columns	<code>SELECT column_name FROM table;</code>
To retrieve unique output records	<code>SELECT DISTINCT column_name FROM table;</code>
For sorting	<code>SELECT col1, col2 FROM table ORDER BY col2;</code>
For sorting backwards	<code>SELECT col1, col2 FROM table ORDER BY col2 DESC;</code>
For counting rows from the table	<code>SELECT COUNT(*) FROM table;</code>
For grouping along with counting	<code>SELECT owner, COUNT(*) FROM table GROUP BY owner;</code>
For selecting maximum values	<code>SELECT owner, COUNT(*) FROM table GROUP BY owner;</code>
Selecting from multiple tables and joining	<code>SELECT pet.name, comment FROM pet JOIN event ON (pet.name = event.name);</code>

### Command Line Statements

Function	Hive Commands
To run the query	<code>hive -e 'select a.col from tab1 a'</code>
To run a query in a silent mode	<code>hive -S -e 'select a.col from tab1 a'</code>
To select hive configuration variables	<code>hive -e 'select a.col from tab1 a' --hiveconf hive.root.logger=DEBUG,console</code>
To use the initialization script	<code>hive -i initialize.sql</code>
To run the non-interactive script	<code>hive -f script.sql</code>
To run script inside the shell	<code>source file_name</code>
To run the list command	<code>dfs -ls /user</code>
To run ls (bash command) from the shell	<code>!ls</code>
To set configuration variables	<code>set mapred.reduce.tasks=32</code>
Tab auto completion	<code>set hive.&lt;TAB&gt;</code>
To display all variables starting with hive	<code>set</code>
To revert all variables	<code>reset</code>
To add jar files to distributed cache	<code>add jar jar_path</code>
To display all the jars in the distributed cache	<code>list jars</code>
To delete jars from the distributed cache	<code>delete jar jar_name</code>

### Metadata Functions and Query

Function	Hive Commands
Selecting a database	<code>USE database;</code>
Listing databases	<code>SHOW DATABASES;</code>
Listing table in a database	<code>SHOW TABLES;</code>
Describing format of a table	<code>DESCRIBE (FORMATTED EXTENDED) table;</code>
Creating a database	<code>CREATE DATABASE db_name;</code>
Dropping a database	<code>DROP DATABASE db_name (CASCADE);</code>