


Zookeeper: Most of it

# ZooKeeper: Why?

- In a distributed system:
  - Servers can crash.
  - Servers can come and go.
  - They must agree on shared information.
  - They must avoid conflicts (e.g., two servers acting as “leader”).
  - They must know who is alive.
- Without a coordinator, they will fight, get confused, or corrupt data
- ZooKeeper ensures everyone behaves properly :)

# ZooKeeper

- ZooKeeper is a small, super-reliable “manager” for distributed systems.
- ZooKeeper is like a class monitor who:
  - Keeps/retains common settings
  - Tells who is alive and dead
  - Decides who will be the leader
  - Coordinates who should do what
  - Sends alerts when something changes.



Data Flair

# ZooKeeper: Core Idea

- Run a small ensemble (3-7) of ZooKeeper servers that replicate a tiny hierarchical data store (znodes)
- Clients read from any server and perform coordinated updates via a leader.
- Leader orders writes with the Zab consensus protocol, giving you consistent configuration, watches, ephemeral and sequential nodes, and primitives.
- You can build locks and leader election on the primitives

## Leader

- .Only one leader at a time.
- .Handles all write requests (create, delete, setData, setACL, multi)
- .Coordinates replication of writes to followers.
- .Runs leader election algorithms.

## **Followers**

- .Usually multiple followers.
- .Handle read requests directly.
- .Forward write requests to the leader.
- .Participate in voting for leader election.

## Observers

- .Optional servers (you may have 0 or many).
- .Handle read requests, just like followers.
- .Do NOT vote in leader election.
- .Do NOT participate in write quorum.
- .Useful for scaling reads without affecting write performance – increase scalability of read operations w/o affecting any writes


## Observers

- .Increase scalability of read operations w/o affecting any writes
- .Reads can come from any node (leader or follower) - Writes must be approved by a majority (quorum) of servers.

If you keep adding MORE servers to increase read capacity, the quorum size becomes larger:

More servers → more majority → writes become slower because the leader must wait for ack from more followers.

Observers are special ZooKeeper servers that: receive updates from the leader, keep a copy of the data, but do not vote in write quorum – so they do not slow down writes



## Znode

- . A ZNode is a node in ZooKeeper's data tree.
- . Each ZNode:
  - . Stores data (small amounts, usually < 1 MB)
  - . Has metadata (version, ACL, timestamps)
  - . Can have children (like directories)

# ZooKeeper: How?

- ZooKeeper keeps a very small tree structure, like folders
- Each item in the tree is called a znode.
- A znode is like a tiny file that stores a few bytes of information.
- Like tiny “notes” for the whole system.



# ZooKeeper: How?

- A node (znode) is a path in ZooKeeper where:
- You can store small data (like a config string)
- You can create children (like directories)
- You can set watches, and
- ZooKeeper keeps it consistent across the cluster.

# Znode examples

- Store app configuration in ZooKeeper
- A microservice starts and registers its presence
- Application tries to acquire a lock.
- A process participates in leader election.
- Prepare the namespace structure before running the app.

# Special Types of znodes

- **Persistent Node**
  - Stays forever until explicitly deleted. /configs/database
  - It does NOT disappear when:
    - The client that created it disconnects
    - The session that created it expires
    - The server restarts (assuming the ensemble is running normally)
  - It stays in ZooKeeper **forever** unless deleted manually or programmatically.

```
zk.create("/myNode", "data".getBytes(),
          ZooDefs.Ids.OPEN_ACL_UNSAFE,
          CreateMode.PERSISTENT);
```

# Special Types of znodes

- **Ephemeral Node**
  - Created by a client
  - Exists only as long as that client's session is alive
  - Disappears automatically when that server dies -
  - Used to know which servers are alive
  - Cannot have children
  - If server-1 crashes, the node auto-deletes → everyone knows it died.

/members/server-1 (ephemeral)

# Special Types of znodes

- **Sequential Node**
  - ZooKeeper adds a number at the end
  - Used for ordering — especially in leader election and locks.


/election/candidate-00001

/election/candidate-00002

# ZooKeeper: Key Components

- **Ensemble** – a group of ZooKeeper servers - availability requires a majority (quorum)
- The group of ZooKeeper servers that work together to form a single, reliable, fault-tolerant ZooKeeper service.
- Think of it as the cluster.

Zookeeper cluster (ensemble)



# ZooKeeper: Key Components

- Leader / Followers / Observers: one of the 3 possible roles played by each server
- The Leader is the only server that:
  - Processes all write requests (create, delete, set data)
  - Assigns transaction IDs (zxids)
  - Proposes updates to all followers
  - Runs ZAB protocol
- Leader is chosen via the ZooKeeper leader election process when the ensemble starts (or when the leader fails).

# ZooKeeper: Key Components

- Leader / Followers / Observers: one of the 3 possible roles played by each server
- Followers are normal voting ZooKeeper servers.
- Handle read requests directly
- Forward write requests to the leader
- Vote on leader proposals (ack)
- Participate in leader election
- Maintain copies of the data tree
- Followers and Leader together form the **quorum** (majority group).
- In 5 node ensemble, if Leader + 2 followers, then ZK operates.

# ZooKeeper: Key Components

- Leader / Followers / Observers: one of the 3 possible roles played by each server
- Observers are **non-voting ZooKeeper servers**.
- Handle read requests
- Forward write requests to the leader
- Do not vote in write agreement
- Do not participate in leader election
- Do not count toward quorum
- Purpose: To scale read-heavy workloads without making the ensemble slower.

# ZooKeeper: Key Components

- **Znodes:** A **znode** is a **node (entry)** in ZooKeeper's hierarchical data tree.
- Store data (up to ~1 MB, but usually a few KB)
- Have children (except ephemeral nodes)
- Be watched by clients
- Be persistent or ephemeral
- Be sequential or not
- Every znode is identified by a **path** like /app/config/version.

# ZooKeeper: Key Components

- Znodes can store:
- Small pieces of data (e.g., "server=active")
- Metadata for distributed systems
- Information for leader election
- Configuration values
- Service discovery information

/brokers/ids/1 (znode for broker 1)

/config/topics/myTopic (topic config)

# ZooKeeper: Key Components

- **Ephemeral nodes** – automatically removed when a client session ends; used for membership & leader election.
- **Sequential nodes** – ZooKeeper appends a monotonically increasing sequence number to the node name (useful for ordered locks/election).

# ZooKeeper: Key Components

- **Watches** – lightweight notification mechanism that let a client know when something changes in the data tree.
- They are one of the most important features of ZooKeeper.
- A watch is a *one-time trigger* that a client sets on a znode.
- When that znode changes, ZooKeeper sends a notification to the client.
- Watches allow clients to avoid constant polling.

# ZooKeeper: Key Components

- Watches allow clients to avoid constant polling.
- Instead of asking repeatedly:
  - "Has the leader changed?"
  - "Has this node been created?"
  - "Has the child list changed?"
- A client places **one watch**, and ZooKeeper notifies it when something happens.

# ZooKeeper: Key Components

- A watch can be triggered by a:
- **Node data change**
- **Node deletion**
- **Node creation**
- **Changes in children list**

set /config "newValue"

delete /leader

exists /election/node-0001 watch      create /election/node-0001 ""

# ZooKeeper: Key Components

- **Zab** (ZooKeeper Atomic Broadcast) — the replication/consensus protocol (leader-based total order broadcast optimized for primary/backup crash-recovery).

# How Leader Election Works

- Every server writes an ephemeral-sequential node:
  - /election/server-00005
- The one with the smallest number becomes leader.
- If leader dies, its ephemeral node disappears → the next smallest number becomes leader.
- So: No fights, no confusion - ZooKeeper ensures fairness.

# Group Membership: Who is alive?

- Each server creates an ephemeral node under /members.
- ZooKeeper keeps the list automatically:

```
/members
  └── server1
  └── server2
    └── server3
```

- If a server crashes → its node disappears → all other servers immediately see the updated list.
- This gives instant cluster membership.

# Watches (Notifications)

- Clients can tell ZooKeeper: “Please tell me if this node changes.”
- This is called a watch.
- When the node changes, ZooKeeper sends a notification.
- Usage:
  - config change alert
  - membership change
  - leader change

# Zookeeper Servers: How?

- ZooKeeper does NOT run as one server: It runs as a small cluster, called an ensemble (usually 3 or 5 servers).
- Why?
  - If 1 server fails, others continue.
  - They vote and agree on updates.
  - It gives very strong reliability.
- ZooKeeper needs a majority to work (called quorum).
- Example:
  - 3 servers → any 2 alive is enough
  - 5 servers → any 3 alive is enough

# Internal Information

- ZooKeeper keeps all data in memory for fast reads.
- When you write something:
  - Write goes to the leader.
  - Leader broadcasts it to followers.
  - Majority confirms.
  - Leader commits the change.
- This ensures the same order of changes everywhere.

# Guarantees

- Total order of updates (a global ordering of state changes).
- Atomicity of updates (multi ops succeed or fail entirely).
- Single system image — clients see an ensemble that behaves like a single datastore.
- Reliability — once a write is acknowledged by quorum it will persist despite failures.
- Simple timeliness / session semantics — ephemeral nodes are tied to client sessions; session timeout determines liveness.

# Common Use cases

- **Leader election** — each contender creates an ephemeral sequential znode under /election; the lowest sequence becomes leader. If it dies, next lowest wins.
- **Group membership / service discovery** — services create ephemeral znodes (e.g., /members/node-00001) so consumers can list current members.
- **Distributed locks** — create ephemeral-sequential nodes under /locks; clients watch predecessor to acquire lock in order.
- **Configuration store** — store small configuration values under znodes; clients watch for changes to react dynamically.
- **Barrier / queue primitives** — implement barriers, work queues using znodes and watches.

# Typical API

- `create(path, data, acl, flags)` — flags: EPHEMERAL, SEQUENTIAL.
- `delete(path, version)`
- `getData(path) / setData(path, data, version)`
- `exists(path, watch)`
- `getChildren(path, watch)`
- `multi(ops)` — atomic transaction of multiple
- `ops.sync(path)` — ensure read is up-to-date with leader.

# ZooKeeper is..

- Not a database (limited data size).
- Not for big data.
- Not for large files.
- Not for high-write workloads.
- It is ONLY for tiny metadata and coordination.

# Dependency on Zookeeper

- Apache Kafka (older versions)
- Hbase
- Hadoop YARN Resource Manager
- Solr
- Storm
- Flink (older versions)

# Observations

- 1) Distributed systems always need some form of coordination
- 2) Programmers cannot use locks correctly
- 3) Message based coordination can be hard to use in some applications

# Wishes

- 1) Simple, Robust, Good Performance
- 2) Tuned for Read dominant workloads
- 3) Familiar models and interfaces
- 4) Wait-Free: A slow/failed client will not interfere with the requests of a fast client
- 5) Need to be able to wait efficiently

# Design Starting Point

Start with the File API and model strip out what we don't need:


- 1)Partial writes/reads (takes with it open/close/seek)
- 2)Rename

add what we do need:

- 1)Ordered updates and strong persistence guarantees
- 2)Conditional updates
- 3)Watches for data changes
- 4)Ephemeral nodes
- 5)Generated file names

# Data Model

- 1) Hierarchical namespace  
(like a file system)
- 2) Each znode has data and children
- 3) data is read and written in its entirety



# ZooKeeper API

String create(path, data, acl, flags)

void delete(path, expectedVersion)

Stat setData(path, data, expectedVersion)

(data, Stat) getData(path, watch)


Stat exists(path, watch)

String[] getChildren(path, watch)


void sync(path)

# Create Flags

- 1)Ephemeral: the znode will be deleted when the session that created it times out or it is explicitly deleted
- 2)Sequence: the path name will have a monotonically increasing counter relative to the parent appended





# ZooKeeper Servers




- 1) All servers store a copy of the data (in memory)
- 2) A leader is elected at startup
- 3) Followers service clients, all updates go through leader
- 4) Update responses are sent when a majority of servers have persisted the change

# HOD

- 1)A client submits a request to start jobtracker and a set of tasktrackers to torque
- 2)The ip address and the ports that the jobtracker will bind to is not known apriori
- 3)The tasktrackers need to find the jobtracker
- 4)The client needs to find the jobtracker




# HOD with ZooKeeper



# HOD with ZooKeeper

When the client spawns the TT and JT tasks in torque, it passes the path of the newly create znode (/hod/jt-1) as a startup parameter.


The client and TT watch the znode for data populated by JT. JT and TT watch the existence of the znode and exit if it goes away.



## Client

```
|  
|-- create znode /hod/jt-1  
|-- spawn JT  
|-- spawn TT  
|  
|-- watches /hod/jt-1 for data from JT  
|  
JT ----> writes data into /hod/jt-1  
    |-- watches /hod/jt-1 existence (exit if deleted)  
TT ----> watches /hod/jt-1 existence (exit if deleted)  
Client -> can delete /hod/jt-1 to signal shutdown
```

# Performance



# Performance at Extremes


Servers	1% Writes	100% Writes
13	265115	4592
9	195178	5550
7	147810	6371
5	75308	8048
3	49827	10519

Numbers are operations per second

# Status

- 1) Project started October 2006
- 2) Prototyped in Fall 2006
- 3) Initial implementation of production service March 2007
- 4) Code moved to zookeeper.sf.net and Apache License November 2007
- 5) Java Quorum and Standalone servers
- 6) Java and C clients available

# ZAB



Why Atomic? Guarantee that each update is applied either completely or not at all across all Zookeeper nodes

Why Broadcast?: Updates from Leader sent to all the followers in the cluster to ensure strong consistency

# ZAB

- ZooKeeper separates how it handles:
  - WRITE requests → complicated, must go through consensus
  - READ requests → very simple, go directly to memory database
- This part is the heart of ZooKeeper's consistency.  
It uses the ZAB protocol (ZooKeeper Atomic Broadcast).
- This module:
  - sends the proposed transaction to all followers
  - waits for ACKs from a majority (quorum)
  - guarantees that all ZooKeeper servers process transactions in the same order
- This ensures strong consistency.

# ZAB

- . ZAB is the consensus protocol used by Apache ZooKeeper to ensure that all ZooKeeper servers (nodes) agree on the same state, even in the presence of failures.
- . Its main goal is to provide highly reliable replication of updates (writes) across a ZooKeeper ensemble (cluster).

# ZAB: How?

- **Leader election** – One node is chosen as the leader.
- **Proposal phase** – Leader receives a client write request, assigns it a unique transaction ID (zxid), and proposes it to followers.
- **Acknowledgment** – Followers persist the proposal and send ACKs to the leader.
- **Commit phase** – When a majority (quorum) of followers ACK, the leader commits the transaction and notifies followers to commit.
- **Recovery after crash** – If the leader crashes, followers use ZAB to elect a new leader and continue from the last committed transaction.

# ZAB: Making it Easy

- . ZAB is a **postal system for distributed updates**.
- . The leader
  - . writes the letter (transaction),
  - . sends copies to all nodes (followers),
  - . waits until enough copies are confirmed (ACKs),
  - . and then officially posts it (commit).
- . Observers just get a copy without participating in the decision.