

## 1) What is HBase? (Start of lecture)

### The simple idea

HBase is a **distributed database** made for **very large data** where you want **fast access**.

Your slide says:

- HBase is a **distributed column-oriented data store** built on top of **HDFS**
- It's an Apache project to provide storage for Hadoop distributed computing
- Data is logically organized into **tables, rows, columns**

### VERY IMPORTANT: Where does it live?

HBase is not independent storage. It stores its files inside HDFS. The diagram says:

- **HBase is built on top of HDFS**
- **HBase files are internally stored in HDFS**

Story:

Think of HDFS as the **warehouse** (file storage).

HBase is the **inventory system** that makes the warehouse searchable and updatable quickly.

---

## 2) HBase is NOT an RDBMS (THIS COMES IN EXAM)

The slide says:

HBase looks like a table, but it is **NOT a relational table**.

So:

- No SQL joins like MySQL (joins aren't supported)
  - No "search any column anytime" style querying
  - HBase is built for **key-based access**
- 

## 3) The HBase Data Model (MOST IMPORTANT CONCEPT)

Your slide gives the most famous HBase definition:

HBase is a **sparse, distributed, multidimensional, sorted map**  
 $(\text{RowKey}, \text{ColumnFamily}, \text{ColumnQualifier}, \text{Timestamp}) \rightarrow \text{Value}$

Let's unpack each word slowly:

### (A) “Map” meaning

Map means like a dictionary:

- input → output  
Here, input is 4-part key, output is the value.

### (B) “Multidimensional” meaning

Because the key has **multiple dimensions**:

1. RowKey
2. ColumnFamily
3. ColumnQualifier
4. Timestamp

### (C) “Sorted” meaning

Sorted primarily by RowKey (lexicographically).

### (D) “Sparse” meaning

HBase stores only what exists.

If a column doesn't exist in a row, no memory/disk is wasted.

**Big exam line:** HBase is best when different rows can have different columns.

---

## 4) Row Key (Your teacher LOVES this)

Slide says:

- Every row is identified by a single unique **row key**
- Data is sorted **lexicographically** by row key (not numeric sorting)
- Range scans are extremely fast due to this sorting
- RowKey is a byte array: `byte[]` (everything converted to bytes)

### Lexicographic sorting (super important trap)

Lexicographic means dictionary-like sorting.

Slide gives example:

If numbers stored as strings:

- "1", "2", "10", "20"
- Sorted becomes:
- "1", "10", "2", "20"

But if stored as fixed-width binary (Bytes.toBytes(int)):

- 1, 2, 10, 20 properly sort numerically.

## What can be a RowKey?

Slide lists: String, Integer, Long, Timestamp, UUID, composite keys, hashes, encoded binary keys — everything becomes byte[].

## Why byte arrays?

Slide answers:

- Flexibility (client can encode any datatype)
- Performance (byte comparison is fast)
- No overhead (HBase doesn't interpret values)
- Efficient lexicographic sorting (regions split easily)

 **Exam tip:** If they ask “why byte[]?”, write these 4 points.

---

## 5) Column Families (second huge pillar)

Slide says:

- HBase groups columns into **column families**
- CFs are fixed at schema definition (can't add/remove frequently)
- Data in the same CF is stored together on disk (same HFile) — **very important**
- Each column belongs to a CF: **cf:column**

## What does “stored together” mean?

Slide explains:

For row user\_001:

- CF personal stored in one HFile
  - CF orders stored in another HFile
  - CF activity stored in another HFile
- You can scan only one CF without touching others.

### **Story:**

Imagine CF is like “folders”:

- personal folder
- orders folder
- activity folder

If you only need “personal”, you don’t open the other folders → faster reads.

## **Creating CFs (syntax in slide)**

You define CFs at table creation:

```
create 'users', 'personal', 'orders', 'activity'
```

## **CF impact on reads/writes (why CF is heavyweight)**

Slide says each CF has its own:

- MemStore
- HFiles
- BlockCache
- Bloom filters

Also:

- Performance tuning is per-CF
- Compaction is CF-wise (big CF → heavy compaction)

## **Column family settings (must know)**

Slide table:

- TTL: auto-delete after X seconds
- Max versions
- Compression (Snappy/GZip/LZ4)
- Block size (default 64KB)
- Bloom filters reduce disk seeks

## **Column family isolation (why split CFs)**

Slide gives example:

- personal: small rarely updated → no compaction pressure
- orders: large frequently updated → separate compaction needed
- activity: high-write time-series → short TTL

## How many column families?

Slide says: **DO NOT make many CFs**

Best practice:

- Max 2–3 CFs
- up to 5 risky
- 1–2 ideal

Why many CFs are bad:

each CF has separate MemStore/HFiles/WAL edits/flushes/compactions → wastes RAM and causes compaction storms.

---

## 6) Column Qualifiers (how HBase becomes “schema-less”)

Within a column family, you can create **unlimited dynamic columns** called qualifiers.

Slide says:

- qualifiers are not predefined; can vary row to row
- gives schema-less flexibility within each family

Another slide adds:

- you can create new qualifiers anytime while inserting (no schema change)
- different rows can have different qualifiers
- HBase stores only what exists → sparse table

Qualifiers are also stored as byte[] and sorted lexicographically.

Qualifiers are “lightweight”:

- don't create new HFiles
- don't create new MemStores
- don't affect WAL structure
- can be millions per row (time-series)

 **Exam line:** Column families are heavyweight; qualifiers are cheap.

---

## 7) Versions + Timestamps (HBase's special power)

This is one of the most testable parts.

### Cell uniqueness formula (again)

(row, cf, qualifier, timestamp) → value

Slide says:

- Versions are not new rows
- they are stored under same row key and same column, with different timestamps

### What is a version?

A version is identified by timestamp.

Latest version returned by default.

### Default versions + changing it

HBase keeps **1 version by default**.

You can change it per CF:

```
alter 'users', NAME=>'info', VERSIONS=>3
```

### Server vs user timestamp

- If you don't provide timestamp → server assigns current time in ms
- You can set your own timestamp (useful for time-series, CDC, event logs, IoT, ordering)

Slides also explain Unix epoch basics (1970 baseline).

### How versions stored in HFiles

Slide says:

- Stored sorted with most recent version first (fast read)
- Versioning helps: analytics at time X, audit, recovery, drift tracking, append-only log style

## **What is a “cell”?**

Smallest unit:

(row key, cf, qualifier, timestamp) → value

---

## **8) HBase internal organization view (logical → internal tree)**

Slide shows the structure:

RowKey → CF → Qualifier → Timestamp → Value

---

## **9) LSM Tree (how HBase achieves fast writes)**

Slides explain HBase uses a Log-Structured Merge Tree.

### **Why LSM?**

Optimized for **high write throughput**.

### **LSM write story**

1. Writes first go to RAM (MemStore/MemTable)
2. Also appended to WAL for durability
3. MemStore keeps sorted in memory
4. When full → flush to disk as immutable file (HFiles/SSTables)
5. Over time many files accumulate → periodic merge/compaction

Slide says it's fast because:

- sequential writes (append WAL, flush)
  - minimized disk seeks
  - reads aided by indexes, Bloom filters, compactions
-

# 10) Major benefits (why HBase is used)

Slide lists:

- sparse (different cols no storage penalty)
  - wide (millions of columns)
  - flexible schema (dynamic qualifiers)
  - high write throughput (LSM)
  - scalable (regions split and spread across region servers)
- 

# 11) Physical model (how HBase stores things in cluster)

Slide says:

**HBase Table → Regions → RegionServers → HFiles**

## (A) Region

A region is a **horizontal partition** of the table based on row-key range.

Example:

- Region1 rows “A” to “H”
- Region2 rows “H” to “Q”

Each region physically has:

- MemStore (RAM)
- Store (disk)
- HFiles
- WAL

## (B) RegionServer

A RegionServer hosts many regions.

Each region has 1 or more stores (one store per column family).

## (C) Store (important)

Column Family = Unit of storage.

For each CF, HBase creates a store with:

- MemStore
  - HFiles on HDFS
- 

## 12) Write path (**THIS IS EXAM GOLD**) — **Put → WAL → MemStore → HFile**

Slides explain this clearly:

When client writes (Put):

1. Write goes to **WAL** (crash recovery)
2. Write goes to **MemStore** (RAM)
3. ACK sent to client
4. When MemStore hits threshold → flush to HFiles on HDFS

MemStore is per column family per region.

### **What is WAL exactly?**

Slide says:

- every write goes to WAL before hitting MemStore
- stored in HDFS as HLog files
- used to replay writes if RegionServer crashes
- prevents data loss

### **Crash scenario (very testable)**

Slide gives scenario:

- inserted 1000 records, still in MemStore
- regionserver crashes  
Without WAL → lost.  
With WAL:
  - master assigns region to another RS
  - new RS reads WAL for that region
  - WAL replay rebuilds MemStore then flushes to HFiles
  - data restored

 If exam asks “why WAL?” write that story.

## **HFile ordering (important detail)**

HFile is sorted:

RowKey → ColumnFamily → ColumnQualifier → Timestamp

---

# **13) Compaction (why reads stay fast over time)**

Motivation slide says:

- frequent small writes → many HFiles in a column family
- reads check all HFiles → slow
- deleted/old versions still occupy space  
Compaction merges and cleans HFiles

### **Minor compaction**

- merges a small number of HFiles
- keeps old versions
- triggered when HFiles exceed threshold (default 3)

### **Major compaction**

- heavier, can impact performance
  - removes tombstones (deleted markers)
  - removes expired versions (TTL/versions)
  - creates a single large HFile per CF per region
- 

# **14) Compression (different from compaction)**

Slide says:

- compression reduces disk usage/I/O by storing data compressed inside HFiles
- codecs: LZO, LZ4, Snappy, GZIP
- applied per column family

- done when writing HFiles (flush)
- does NOT merge files or remove old versions

And it gives comparison table (ratio vs speed vs CPU).

---

## 15) Cache (BlockCache + MemStore) + eviction policy

Slide says:

- Block Cache stores hot data blocks read from HFiles (64KB blocks)
  - MemStore caches data written but not flushed
  - eviction policy: LRU (least recently used)
- 

## 16) HBase Architecture components (Master, RegionServer, Client)

Slide says 3 major components:

- HBaseMaster (one master)
- HRegionServer (many region servers)
- HBase client

Master responsibilities:

- coordinates slaves
- assigns regions, detects failures
- admin functions

RegionServer:

- manages regions
  - serves reads/writes
  - uses log (WAL)
-

# 17) HBase vs HDFS (where each is good)

Slide says:

HDFS:

- good for batch processing (scans over big files)
- not good for record lookup, small incremental adds, updates

HBase:

- excellent for fast random reads by row key (ms lookup)
- high throughput for large datasets

Not good for:

- ad-hoc lookups by non-key columns
- secondary indexes (not built in)
- random scans over large ranges (slow/expensive)

**One clean line from slide:**

HBase is good for key-value pattern lookups, not “query anything anywhere.”

---

# 18) Write strategies (Random, Bulk load, Incremental)

Slide says:

## Random writes

HBase is well suited for random writes due to LSM:

MemStore → flush to HFiles.

## Bulk load

For huge ingestion:

- prepare sorted HFiles offline
- directly load into regions
- efficient for ETL pipelines

## Incremental load

Small updates in real-time:

- events streamed incrementally
- each write results in a new version

Also slide warns: small writes cause RPC overhead, frequent flushes, many tiny HFiles, compaction overhead → better buffer updates and write in bulk.

---

## 19) Read strategy (Get, Scan + caches + bloom filters)

Slide says:

- **Get** retrieves a single row by row key
- **Scan** retrieves multiple rows in a row-key range

Block cache:

- check cache first before disk

Bloom filters:

- quickly determine if row/column exists in HFile without reading entire file (reduces disk seeks)
- 

## 20) Operations (Put / Delete) + Versions

Slide says:

- Put can insert new record or update existing key
- implicit timestamp (server) or explicit timestamp

Delete:

- marks cells as deleted (tombstones)
- can delete at multiple levels: entire CF, all CFs of row, etc.

Joins:

- HBase does not support joins
  - can be done in application layer using get/scan
- 

## **Exam-ready “3 point” answers (most likely questions)**

### **Q1) Why HBase uses RowKey sorting?**

1. Rows sorted lexicographically by RowKey
2. Fast point lookup using exact key
3. Range scans become extremely fast due to sorting

### **Q2) Explain Column Family vs Column Qualifier**

- CFs are fixed in schema and stored together in one HFile (heavyweight).
- Qualifiers are dynamic, schema-less, cheap, can be millions.
- CF design affects compaction/caching/performance.

### **Q3) Explain write path (Put)**

1. Client sends Put to RegionServer → append to WAL
2. Write stored in MemStore (RAM)
3. Later flush MemStore → HFile on HDFS; WAL enables replay on crash

### **Q4) Minor vs Major Compaction**

- Minor: merge few HFiles, reduce file count, keeps old versions
- Major: heavy, merges all, removes tombstones + expired versions (TTL/versions)
- Needed because frequent writes create many small HFiles and slow reads