

# 1 FIRST: WHY Lambda Architecture exists (the STORY)

Imagine a system like:

- Uber rides
- Bank transactions
- University attendance
- IoT sensors
- Website clicks

**Problem:**

Data is coming **ALL THE TIME.**

You want:

- **Real-time answers** → “What’s happening RIGHT NOW?”
- **Accurate answers** → “What is the CORRECT total after cleaning errors?”
- **Scalability** → data is huge
- **Fault tolerance** → machines can fail

But here’s the catch:

✗ Real-time systems are **fast but sometimes wrong/incomplete**  
✗ Batch systems are **accurate but slow**

👉 Lambda Architecture exists to get **BOTH speed + accuracy**

---

# 2 WHAT IS LAMBDA ARCHITECTURE (EXAM DEFINITION)

Lambda Architecture is a big data architecture that processes data using two parallel paths:

- a **batch layer** for accurate historical computation

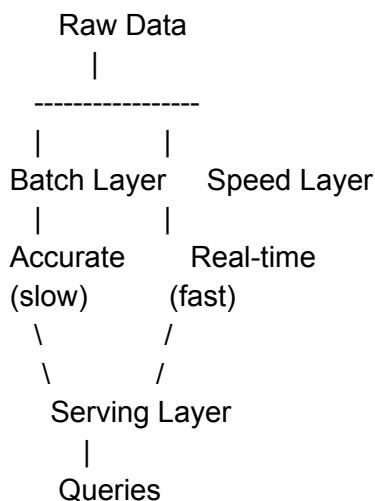
- a **speed layer** for real-time processing and merges results in a **serving layer**.

## **Keywords your teacher wants:**

- batch processing
  - real-time / stream processing
  - fault tolerance
  - scalability
  - low latency + correctness

# **3 THE CORE IDEA (VERY IMPORTANT)**

👉 Same data goes to TWO paths at the same time



You are NOT choosing batch OR streaming.  
You are doing **BOTH** in parallel.

## **4 THE THREE LAYERS — IN DEPTH (THIS IS EXAM GOLD)**

## 1. BATCH LAYER (the “truth keeper”)

### What it does

- Stores **ALL historical data**
- Recomputes results **from scratch**
- Gives **accurate, correct results**

### Characteristics

- Slow (minutes / hours)
- Very accurate
- Fault-tolerant (can recompute)
- Uses large-scale distributed processing

### Technologies (Databricks side)

In exam, you write:

- **Databricks (Spark Batch)**
- Storage: **HDFS / Data Lake / S3 / DBFS**
- Processing: **Spark jobs on Databricks**

### Example

“Every night, recompute total sales per product from ALL orders since day 1.”

👉 If there was bad data earlier, batch layer **fixes it**.

---

## 2. SPEED LAYER (the “fast but rough” layer)

### What it does

- Processes **new incoming data immediately**
- Handles data **not yet processed by batch**
- Gives **low-latency results**

### Characteristics

- Very fast (seconds)
- Might be incomplete
- Uses streaming

- Temporary results

## Technologies (Databricks side)

In exam:

- **Databricks Structured Streaming**
- Input: **Kafka / Event Stream**
- Processing: Spark Streaming on Databricks

## Example

“Show orders from the last 30 seconds.”

👉 This layer **fills the gap** while batch is running.

---



## 3. SERVING LAYER (the “answer machine”)

### What it does

- Combines **batch results + speed results**
- Serves queries to users/apps
- Optimized for fast reads

### Characteristics

- No heavy computation
- Read-optimized
- Always query-ready

## Technologies

In exam:

- **HBase / Cassandra / MongoDB**
- Or Databricks Delta tables queried by BI tools

## Example

Dashboard asks:

“Total sales till now”

Serving layer answers using:

batch\_result + speed\_result

---

## 5 WHERE DATABRICKS FITS (THIS IS CRITICAL)

Databricks is **NOT** a single layer.

It can be used in **MULTIPLE** layers.

**Databricks roles:**

Layer	Databricks Role
Batch Layer	Spark batch jobs
Speed Layer	Structured Streaming
Serving Layer	Delta Tables / SQL / BI

👉 This is WHY Databricks is perfect for Lambda Architecture.

---

## 6 EXAM-READY ARCHITECTURE #1 (CLASSIC LAMBDA)

**Use this when you want SAFE MARKS**

**Scenario:** E-commerce analytics

**Flow:**

1. Events (orders/clicks) generated
2. Events go to **Kafka**
3. Kafka sends data to:
  - Databricks Batch
  - Databricks Streaming

4. Batch computes historical aggregates
5. Streaming computes real-time aggregates
6. Results stored in **HBase**
7. Dashboard queries HBase

**Write in exam:**

- Batch: Databricks Spark batch
  - Speed: Databricks Structured Streaming
  - Serving: HBase
- 

## 7 EXAM-READY ARCHITECTURE #2 (DELTA-BASED LAMBDA)

This is **very Databricks-specific** and looks advanced.

**Flow:**

1. Data ingested into **Delta Lake**
2. Batch jobs recompute tables nightly
3. Streaming jobs update same Delta tables
4. Delta handles versioning + ACID
5. BI tools query Delta tables

**Why examiner likes this:**

- Shows Databricks expertise
  - Mentions Delta Lake
  - Shows simplification of serving layer
- 

## 8 EXAM-READY ARCHITECTURE #3 (IOT / SENSOR SYSTEM)

**Scenario:** Smart campus / sensors

**Batch Layer**

- Databricks batch jobs analyze historical sensor data

### Speed Layer

- Streaming detects abnormal temperature / motion

### Serving Layer

- Alerts + dashboards

#### Write example:

“Speed layer triggers alerts instantly, batch layer corrects long-term trends.”

---

## 9 EXAM-READY ARCHITECTURE #4 (BANKING / FRAUD)

#### Scenario: Fraud detection

- Speed layer flags suspicious transaction immediately
- Batch layer recalculates risk scores overnight
- Serving layer merges both

#### Key exam phrase:

“Speed layer provides immediate reaction, batch layer ensures correctness.”

---

## 10 HOW TO DRAW THIS IN THE EXAM (STEP-BY-STEP)

#### Always draw:

1. **Data source at top**
2. Split arrow into TWO paths
3. Label:
  - Batch Layer (Databricks Spark)
  - Speed Layer (Databricks Streaming)

4. Merge into Serving Layer
5. Arrow to Users / Dashboard

👉 Even a **simple clean diagram** gets marks.

---

## **11 WHAT TO WRITE IN WORDS (20-MARK STRUCTURE)**

Use this structure:

- 1. Definition of Lambda Architecture (2–3 lines)**
  - 2. Explain Batch Layer (role + tech)**
  - 3. Explain Speed Layer (role + tech)**
  - 4. Explain Serving Layer (role + tech)**
  - 5. Databricks integration**
  - 6. Real-world example**
  - 7. Advantages**
  - 8. Limitation (complexity)**
- 

## **12 COMMON EXAM TRAPS (AVOID THESE)**

- ✗ Saying Lambda = only streaming
- ✗ Forgetting serving layer
- ✗ Not mentioning Databricks role
- ✗ Not giving example
- ✗ No diagram

---

13

## ONE-PAGE MEMORY VERSION (LAST-MINUTE)

- Lambda = Batch + Speed + Serving
- Batch = slow + accurate + Databricks Spark
- Speed = fast + streaming + Databricks
- Serving = query layer
- Same data flows to BOTH
- Databricks can do BOTH batch + streaming