

## 0) What MongoDB actually is (foundation in 60 seconds)

MongoDB is a **document database**.

- Data is stored as **documents** (like JSON objects), grouped in **collections**.
- Internally MongoDB stores documents as **BSON (Binary JSON)**, which is faster and supports more types than plain JSON.

### JSON vs BSON (why BSON matters)

- **JSON** = human readable, used in web APIs.
- **BSON** = binary format Mongo uses internally for speed + richer type support.

### Biggest Mongo idea: “Embed instead of Join”

In SQL, you often **JOIN** tables. Mongo prefers **embedding** related data inside the same document to avoid joins.

Example from slide idea:

- SQL join: users + addresses
- Mongo: user document contains an array of addresses (embedded).

**Exam line:** “MongoDB reduces joins by embedding documents; referencing exists but no enforced foreign keys like SQL.”

---

## 1) Data modeling (embedding vs referencing) — this is asked a LOT

### A) Embedding (put child inside parent)

Use when:

- You usually read the data together
- Child data belongs mostly to one parent
- Child list is not exploding infinitely

Example idea (books + authors embedded):

A book document contains authors array; query can search inside array fields.

## B) Referencing (store IDs, like foreign keys)

Use when:

- Many-to-many
- Data reused by many parents
- Data changes frequently and you don't want duplication

**Exam trap:** Mongo allows referencing, but it **doesn't enforce** foreign keys.

### Modeling scenario (very exam-likely)

“Book checkout system”

- One book can be checked out by one student at a time
- Student can check out many books

A clean Mongo approach:

- Put `checked_out` array inside student (embed checkout history), each item stores `book_id + date`.

#### Exam answer structure (3 points):

1. Why embed? faster reads, fewer joins
  2. What's the risk? duplication / large document
  3. When reference? many-to-many / shared entities
- 

## 2) Indexing (highly likely short question)

### What happens without index?

Mongo scans **every document** → slow.

### With index

Mongo uses an index to jump directly to matching docs.

Your lecture lists what they can ask:

- single field index
- compound index
- multikey index

- `explain`
- compare with/without index

#### Exam-friendly points (memorize):

- Index = extra data structure that speeds reads
  - Tradeoff = more storage + slower writes (because index updates)
  - Use `explain()` to see if query used index and how efficient
- 

## 3) Aggregation vs MapReduce (common “compare” question)

### Aggregation pipeline (preferred)

Aggregation = operations that process records and return results.

Pipeline provides:

- filters like queries
- transformations for output
- grouping/sorting/array aggregation/operators

Examples from slides (understand meaning, no need to memorize full code):

- `$group` for totals/average
- `$match` filter then `$project` show selected fields
- `$sort` then `$limit` top k

### MapReduce (older / heavier)

- map stage emits intermediate key-values
- reduce stage combines results
- optional finalize stage
- uses custom JavaScript functions, flexible but less efficient/complex
- can output bigger results than pipeline 16MB limit

#### Exam compare (3 points):

1. Pipeline simpler + faster; MapReduce more flexible but heavy
  2. Pipeline is optimized inside MongoDB; MapReduce uses JS functions
  3. MapReduce handles bigger outputs; pipeline has 16MB output limit
-

# NOW THE MAIN EXAM PART: “Distributed MongoDB”

Your guideline says this is **most important**: primary/secondary, reading/writing, oplog, read/write concern, arbiter, sharding, load balancing.

## 4) Replication = Replica Set (high availability + redundancy)

### What is replication?

A **replica set** = group of MongoDB servers that maintain the same data.

- **One Primary**: handles writes (and reads by default)
- **Secondaries**: replicate data from primary; can serve reads if allowed

### Primary node job

- accepts writes
- sends changes to secondaries using **oplog**

### Secondary node job

- replicates from primary using oplog
  - can serve reads if read preference allows
  - cannot accept writes unless promoted to primary
- 

## 5) Oplog (the “log” that makes replication possible)

**Oplog** = operations log: an ordered record of changes on the primary.

Story:

1. Primary receives write (insert/update/delete)
2. Primary records it in oplog
3. Secondaries “tail” the oplog and apply the same operations

**Exam line:** “Secondaries replicate data from the primary using oplog.”

---

## 6) Automatic failover + Elections (very exam-likely scenario)

### What happens when primary dies?

Slides give the flow:

- Secondaries detect failure using **heartbeats** (every 2 seconds)
- Election starts (Raft-like consensus)
- One secondary becomes new primary
- Clients auto-reconnect
- downtime about **2–5 seconds**

Extra election knobs (important terms):

- election timeout default 10s (`electionTimeoutMillis`)
- priority affects who becomes primary
- arbiters help maintain odd votes

#### Exam scenario answer (write this style):

"When primary fails, secondaries notice via heartbeats, trigger election, choose new primary based on voting/priority, and clients reconnect automatically."

---

## 7) Arbiter (aik aur node) — what it is and why we use it

Arbiter:

- participates in voting (helps quorum)
- **does not store data**
- **cannot become primary**
- used when you have even number of data-bearing nodes to avoid tie

#### Exam 3 points:

1. Adds a vote → helps elect a primary (quorum)
  2. No data storage → cheaper, but not a backup
  3. Cannot be primary → only decision-maker, not worker
-

## 8) Read preference vs Read concern (don't mix!)

These sound similar but are different.

### A) Read preference = “From where do I read?”

- Default reads are from **primary** (stronger consistency)
- You *can* read from secondary if read preference allows

### B) Read concern = “How fresh/committed must the data be?”

From Lecture 15:

- **local** (default) may see uncommitted
- **available** low consistency
- **majority** only majority-committed data
- **linearizable** strongest, slowest

#### Story example (very exam):

If you read from a secondary right after a write, that secondary may not have caught up (replication lag). Using stronger read concern like **majority** ensures you only see data confirmed by majority nodes, but it may be slower.

---

## 9) Write concern = “How many nodes must confirm my write?”

From lecture:

- **w:1** → only primary confirms (fast, weaker safety)
- **w:majority** → most nodes confirm (safer)
- **w:0** → fire-and-forget (fastest, risky)

#### Exam scenario: Bank transfer

- Use **w:majority** because you cannot lose a confirmed write if primary crashes right after acknowledging.
- 

## 10) Mongo consistency model (CP-like vs AP-like)

Lecture 15 says Mongo default:

- Writes: primary only (`w:1`)
- Reads: primary (`local`)

And you can tune it:

- CP-like (stronger consistency)
- AP-like (more availability)
- tradeoff: better consistency = slower; better performance = weaker guarantees

**Exam one-liner:** “MongoDB can be tuned along the consistency–availability tradeoff using read/write concerns and read preference.”

---

## 11) Sharding (the BIG second half)

### What is sharding?

Sharding = **horizontal scaling**: split one big dataset across many machines.

Why?

When workloads become slow if:

- data > RAM
  - data > single machine disk
  - high write/transaction rate
  - want linear scalability
- 

### Sharded cluster components (must know diagram words)

#### 1) `mongos` (query router)

- lightweight router
- receives client queries
- routes to correct shard(s)

#### 2) Config servers (CSRS)

- store cluster metadata (chunk locations)
- must maintain quorum; without quorum metadata access stops
- three-node replica set required

### 3) Shards

- each shard is itself a **replica set**
  - gives scalability + availability
- 

## 12) Shard key (**super important exam topic**)

Every sharded collection chooses a **shard key**; it decides which shard stores which document.

### Good shard key properties (memorize)

- high cardinality (many unique)
- evenly distributed
- frequently used in queries
- not monotonically increasing (or use hashed)

Bad keys:

- timestamp (monotonic → hotspot)
- boolean (only two values)
- small-category field

### Hotspotting (term sir LOVES)

Hotspotting = too many reads/writes hit same shard/chunk → that shard overloaded, others idle  
→ performance imbalance.

---

## 13) Chunks + Balancer = load balancing

Mongo breaks data into **chunks** (default 64MB).  
Chunks can move between shards based on load.

Balancer:

- background process
- monitors chunk distribution

- moves chunks when uneven
- migrations happen online while cluster stays active

**Exam line:** “Sharding scales; balancer ensures even distribution by moving chunks online.”

---

## 14) Read/write flow in sharded cluster (the scenario question)

### Write flow

Client → `mongos` → correct shard → primary of that shard

### Read flow depends on query

If query includes shard key → goes to only one shard (fast)

If query doesn't include shard key → broadcast to all shards (**scatter-gather**, slow)

**Exam tip sentence:** “A good shard key avoids scatter-gather; otherwise queries hit all shards and become slow.”

---

## 15) Failover inside sharded clusters (short but important)

If one shard's primary fails:

- that replica set elects new primary
- only that shard temporarily degraded
- other shards continue normally

If config servers fail:

- must maintain quorum
- without quorum cluster metadata access stops

---

## 16) Distributed transactions (advanced but examinable)

Mongo supports multi-document ACID transactions across shards (v4.2).  
Works like **two-phase commit (2PC)** to ensure atomicity across docs/shards.  
But it's slower; avoid unless necessary.

Also mentioned:

- oplog coordination
- retryable writes (safe retry after network error/primary failover)

**Exam scenario:**

"Updating two collections on different shards must be atomic" → use transactions (2PC), but mention performance tradeoff.

---

## 17) Backup + monitoring (easy marks)

Backup options:

- `mongodump/mongorestore` (small DB)
- filesystem snapshots
- Cloud Manager / Ops Manager  
For sharded cluster: backup each shard + config server

Monitoring tools:

- mongostat, mongotop
- Atlas metrics, Ops Manager dashboards  
Metrics:
  - replication lag
  - chunk imbalance
  - page faults
  - qps, connections, cache usage

---

## High-yield “Sir-style” questions (with perfect answer skeletons)

**Q1) Explain replica set and role of primary/secondary + oplog.**

**Answer skeleton (3 points + example):**

- Replica set maintains same data for availability; one primary handles writes; secondaries replicate
- Primary records changes in **oplog** and secondaries apply them
- Example: if primary crashes, secondary becomes new primary via election → service continues

## **Q2) What is arbiter and why used?**

- Votes in election (quorum)
- No data, cannot be primary
- Used to avoid tie (odd votes)

## **Q3) Define write concern and compare w:1 vs w:majority.**

- Controls #nodes that confirm write
  - w:1 fast but weaker
  - w:majority safer but slower
- Example: banking uses majority.

## **Q4) Define read concern and compare local vs majority vs linearizable.**

Use meanings from slide, mention tradeoff.

## **Q5) What is sharding and why needed?**

Mention scale-out, data too big, high writes, linear scalability.

## **Q6) Explain sharded cluster architecture (mongos, config servers, shards).**

Use exact roles.

## **Q7) What makes a good shard key? What is hotspotting?**

Give properties + define hotspotting.

## **Q8) Explain scatter-gather and why it is slow.**

When query lacks shard key, it broadcasts to all shards.

---

## **5-minute cram list (if time is dying)**

- Replica set: primary writes, secondary replicates with **oplog**

- Failover: heartbeats → election → new primary in seconds
- Arbiter: vote only, no data
- Write concern: w:1 vs w:majority
- Read concern: local/majority/linearizable
- Sharding: mongos + config servers + shards (replica sets)
- Shard key: high cardinality, even distribution; avoid hotspotting
- Scatter-gather slow without shard key
- Chunks + balancer moves data online