

## 1 First: WHY Kafka exists (the story)

Imagine this situation:

- A website is generating **thousands of events per second**
  - user clicks
  - orders
  - payments
  - logs

Now the problem:

- Producer (website) is **fast**
- Consumer (analytics / billing) is **slow**
- If consumer crashes → data lost
- If producer waits → website slows down

### ✗ Direct communication fails

Producer → Consumer directly = **tight coupling**, failures, data loss.

### ✓ Solution: Kafka in the middle

Kafka acts like a **buffer + log + pipeline**.

👉 Producer throws data into Kafka and forgets

👉 Consumer reads whenever it is ready

This is the core idea of Kafka.

---

## 2 Definition (WRITE THIS IN EXAM)

Apache Kafka is an open-source distributed event streaming platform designed for high-throughput, fault-tolerant, real-time data pipelines and stream processing.

Key words your teacher wants:

- distributed
- event streaming
- real-time

- fault tolerant
  - high throughput
- 

## ③ Kafka basic building blocks (VERY IMPORTANT)

### (A) Topic

A **topic** is a **named stream of records**

Think of it like:

- a table name
- or a queue name

Examples:

- `orders`
- `clickstream`
- `user_logs`

👉 Producers write to topics

👉 Consumers read from topics

---

### (B) Partitions (MOST IMPORTANT)

Each topic is split into **partitions**.

Each partition is:

- **Ordered** → messages have strict sequence
- **Immutable** → once written, never changed
- **Parallelizable** → multiple consumers can read in parallel

So:

Topic: `orders`

Partitions: P0, P1, P2

---

### (C) Offset (EXAM FAVORITE)

Inside **each partition**, every message has an **offset**.

**Offset = position of message inside a partition**

- Offsets start from  $0$
- Increase sequentially
- Unique **only within that partition**

👉 Very important:

Kafka **does NOT track consumption — consumer does** using offsets.

---

## 4 Producer (who sends data)

A **producer**:

- Sends messages to Kafka
- Chooses **which partition** to send to

**How partition is chosen?**

- If **message key is present**  $\rightarrow \text{hash}(\text{key}) \% \text{ num\_partitions}$
- Same key  $\rightarrow$  **same partition** (order preserved)

Example:

key = "Lahore"  $\rightarrow$  always P0

key = "Karachi"  $\rightarrow$  always P2

👉 This guarantees **ordering for same key**

---

## 5 Broker & Kafka Cluster

**Broker**

A **broker** is a **Kafka server**.

- Stores messages
- Serves producers & consumers
- A cluster has **multiple brokers**

Each broker can host:

- multiple partitions
  - from multiple topics
- 

## 6 Replication & Fault Tolerance (VERY IMPORTANT)

Each partition is **replicated**.

- **Leader replica** → handles read & write
- **Follower replicas** → copy data

If leader fails:

👉 one follower becomes **new leader** automatically

### ISR (In-Sync Replicas)

ISR = replicas fully caught up with leader.

Kafka commits a write **only when ISR confirms**.

👉 This ensures **fault tolerance**

---

## 7 Consumer (who reads data)

A **consumer**:

- Reads messages from partitions
- Tracks **offsets**
- Commits offsets after processing

Kafka does **NOT push** data — consumer **pulls** data.

---

## 8 Consumer Group (EXAM GOLD)

Consumers work in **groups**.

Rules:

- Each partition → **only ONE consumer in a group**
- Consumers share load
- Kafka balances partitions automatically

Example:

Topic: orders (3 partitions)

Consumer group: billing-group

Consumers: C1, C2

P0 → C1

P1 → C2

P2 → C2

---

## 9 Offset Commit & Delivery Guarantees

Kafka supports:

### At-most-once

- Message may be lost
- No duplicates

### At-least-once

- Message may be processed twice
- No loss (default)

### Exactly-once

- Processed once
  - Needs special configuration
- 

## 10 Consumer Lag (VERY VERY IMPORTANT)

$$\text{Consumer Lag} = \text{Latest Offset} - \text{Committed Offset}$$

It shows:

- how far consumer is behind producer

Low lag = healthy

High lag = problem

Causes:

- producer too fast
  - consumer slow
  - rebalancing
  - crash/network issues
- 

## 11 Zookeeper (linked with Kafka – EXAM)

Kafka (legacy) uses **Zookeeper** for:

- metadata management
- broker coordination
- leader election

👉 You will study Zookeeper in **later lecture**, but here know:  
Kafka depends on Zookeeper for **cluster coordination**.

---

## 12 Kafka is like a Message Queue (but better)

Kafka provides:

- asynchronous communication
- load leveling
- fault tolerance
- scalability
- loose coupling

Difference:

- Kafka **does not delete messages immediately**
  - Consumers track offsets themselves
- 

## 13 Where Kafka fits in Big Data Architecture

Kafka is used for:

- real-time ingestion
  - streaming pipelines
  - lambda architecture **speed layer**
- 

## EXAM-READY SHORT ANSWERS (3 points)

### **Q1: What is a Kafka topic?**

- A named stream of records
- Divided into partitions
- Used by producers and consumers

### **Q2: Why partitions are important?**

- Enable parallelism
- Preserve ordering within partition
- Improve scalability

### **Q3: What is offset?**

- Position of message in partition
- Used by consumer to track progress
- Enables replay of data

### **Q4: Explain consumer group.**

- Group of consumers
- Each partition assigned to one consumer
- Enables load balancing

## **1) Define Apache Kafka. Mention 3 keywords that must appear in a perfect definition.**

**Answer:** Kafka is an **open-source distributed event streaming platform** for **high-performance real-time pipelines** and **stream processing**. Keywords: **distributed, event**

**stream, fault-tolerant/durable.**

**Example:** capturing website clickstream events and sending them to analytics in real time.

## **2) Kafka vs “traditional message queue”: give 3 differences.**

**Answer:** (1) Kafka stores messages as an **append-only log** (records are immutable) rather than deleting immediately after consumption. (2) Kafka supports **replay** using offsets; queues usually delete after ack. (3) Kafka is designed for **high throughput** + partitioned scalability; queues focus on task distribution.

**Example:** replay all orders from yesterday for debugging using offsets.

## **3) Define “Topic” and give 2 analogies from the slide.**

**Answer:** A topic is a **named stream of records**. Analogies: like a **database table name** or a **queue name**.

**Example:** topic = `orders, clickstream, user_logs`.

## **4) What is a “Partition” and why does Kafka split topics into partitions? (3 reasons)**

**Answer:** A partition is a **sub-log** of a topic (numbered 0,1,2...). Reasons: (1) **parallelism** (multiple consumers read in parallel), (2) **scalability** (spread load across brokers), (3) preserves **ordering within a partition**.

**Example:** `orders` topic with 3 partitions lets 3 consumers process simultaneously.

## **5) Explain “ordered” and “immutable” partitions. Why are these properties important?**

**Answer:** Ordered means messages in a partition have a **strict sequence**. Immutable means once written, messages are **not modified**. Importance: (1) enables correct event history, (2) supports replay, (3) makes replication simpler and consistent.

**Example:** order events must be processed in order for a single user session.

## **6) What is an “offset”? Give 3 properties.**

**Answer:** Offset is a **unique numeric position** of a record inside a partition. Properties: (1) increases monotonically, (2) only meaningful **within that partition**, (3) consumers track

progress by storing committed offsets.

**Example:** consumer processed up to offset 200 ⇒ next read starts at 201.

## 7) Producer vs Consumer vs Broker: define each in one clean paragraph (3 lines each).

**Answer:** **Producer** writes messages to Kafka topics/partitions. **Consumer** reads messages and processes them. **Broker** is a Kafka server that stores/serves messages; a cluster has multiple brokers and each can host multiple partitions.

**Example:** Order Service (producer) → Kafka (brokers) → Billing Service (consumer).

## 8) What is a Kafka “cluster” and why do we need multiple brokers? (3 points)

**Answer:** A cluster is a group of brokers working together. Multiple brokers provide: (1) **horizontal scaling** (more partitions/storage), (2) **fault tolerance** (replicas on different brokers), (3) better throughput by distributing partitions.

**Example:** if broker-1 fails, partitions replicated on broker-2 still survive.

## 9) What is a “consumer group”? Explain with 3 rules.

**Answer:** Consumers can work as a **group**. Rules: (1) partitions are **divided among consumers** for load balancing, (2) **each message goes to only one consumer** in the group, (3) if a consumer dies, partitions are reassigned (rebalance).

**Example:** Analytics group has 3 consumers; each handles some partitions of **clickstream**.

## 10) Scenario: Topic has 3 partitions and consumer group has 2 consumers. What happens?

**Answer:** Kafka assigns partitions across consumers: one consumer may get 1 partition, the other gets 2 partitions. Key points: (1) **one partition cannot be read by two consumers in same group at same time**, (2) load might be uneven, (3) scaling consumers beyond partitions gives no extra parallelism.

**Example:** Consumer-A: P0, Consumer-B: P1 & P2.

## 11) Scenario: Group has 10 consumers but topic has 3 partitions. Explain the outcome.

**Answer:** Only **3 consumers actively read** (one per partition max). Remaining consumers are **idle** in that group. This shows: (1) parallelism limited by partitions, (2) you must plan partitions for scaling, (3) too many consumers wastes resources.

**Example:** for higher throughput, increase partitions instead of adding consumers.

## 12) Explain replication in Kafka: leader replica vs follower replica (3 differences).

**Answer:** Each partition has replicas. **Leader** handles **all reads/writes**. **Followers** copy from leader. If leader fails, a follower can become new leader. Differences: (1) leader serves clients, (2) followers replicate, (3) failover promotes follower to leader.

## 13) What is ISR (In-Sync Replicas) and why does it matter?

**Answer:** ISR is the set of replicas **fully caught up with the leader**. Matters because: (1) defines “safe” replicas, (2) affects durability—writes are considered committed when ISR confirms (as per slides), (3) if replica falls behind it leaves ISR, reducing fault tolerance.

**Example:** if ISR size shrinks to 1, broker failure can risk data availability.

## 14) Explain “commit” in Kafka in 3 points (don’t confuse commit offset vs commit write).

**Answer:** Two “commits”: (1) **write commit**: message considered committed after required replicas (ISR) have it, (2) **offset commit**: consumer saves last processed offset, (3) together they define reliability: durable storage + tracked progress.

**Example:** consumer commits offset 120 after billing done so it won’t reprocess earlier messages.

## 15) Define “at-most-once”, “at-least-once”, “exactly-once” delivery. Also mention risk.

**Answer:** **At-most-once**: may lose messages (data loss). **At-least-once**: no loss but duplicates possible. **Exactly-once**: delivered once only; requires configuration/transactions.

**Example:** payments need exactly-once; logs can tolerate at-least-once with dedup.

## 16) Scenario: Why can at-least-once create duplicates? Give 3 reasons.

**Answer:** Duplicates occur if: (1) consumer processes message but crashes **before committing offset**, (2) producer retries send on timeout, (3) rebalance happens while processing.

**Example:** email service may send same email twice unless dedup by message id.

## 17) Explain how message keys decide partitions. Give the rule and benefit.

**Answer:** Partition selection often uses **hash(key)**. Rule: messages with the **same key go to the same partition**. Benefit: (1) preserves ordering for that key, (2) enables consistent processing, (3) better locality for aggregation by key.

**Example:** key=userA  $\Rightarrow$  all userA events go to P0 (M1 then M4).

## 18) Scenario: When should a producer use a key vs no key?

**Answer:** Use key when: (1) you need ordering per entity, (2) you want same-entity aggregation, (3) you want stable partitioning. No key when: (1) pure load balancing is more important.

**Example:** key by `customer_id` for banking events; no key for random metrics logs.

## 19) Define consumer lag. Provide the exact formula.

**Answer:** Consumer lag is how far behind a consumer is from the latest offset. Formula:  
**Lag = Latest Offset (log end) – Last Committed Offset.**

**Example:** latest=4, committed=2  $\Rightarrow$  lag=2 (two messages waiting).

## 20) Numerical: Latest offset=1200, consumer committed=1188. Compute lag and interpret.

**Answer:** Lag =  $1200 - 1188 = 12$ . Interpretation: consumer is **12 messages behind** in that partition; indicates slower consumption or temporary issues.

**Example:** analytics dashboard freshness may drop until lag returns near 0.

## 21) Numerical (cluster lag): Partition lags are 5, 0, 11, 2. What is cluster lag and meaning?

**Answer:** Cluster lag = sum = **18**. Meaning: across partitions, consumers are behind by total 18 messages; the system may face delays/SLAs issues if it grows.

## **22) What causes consumer lag? Give at least 4 reasons from lecture.**

**Answer:** Causes: (1) producer rate too high, (2) consumer processing slow/heavy, (3) consumer failure/disconnection, (4) rebalancing partition reassignment, (5) network/disk issues affecting reads/offsets.

**Example:** if consumer does ML inference per message, it may fall behind.

## **23) Low lag vs high lag: explain impact in 3 points.**

**Answer:** Low lag means consumer keeps up and data is fresh. High lag means falling behind → delays, potential timeouts, stale dashboards. High lag can also signal under-provisioning or a bug in consumer.

**Example:** fraud detection lagging by minutes can miss real-time prevention.

## **24) Explain Kafka “data flow” step-by-step (4 steps).**

**Answer:** (1) Producer sends event to broker/topic partition, (2) broker writes to disk and replicates to followers, (3) consumer reads from leader partition and tracks offset, (4) consumer commits offset to ensure at-least-once delivery.

**Example:** order placed → stored in Kafka → billing reads → commits offset.

## **25) Explain “load balancing” inside consumer groups (3 points).**

**Answer:** (1) Kafka assigns partitions among consumers, (2) each partition goes to one consumer in that group, (3) when consumers join/leave, rebalance reassigns partitions.

**Example:** adding Consumer-C can redistribute partitions to reduce lag.

## **26) What is “rebalance” and why can it create lag spikes?**

**Answer:** Rebalance is partition reassignment across consumers. It can spike lag because: (1) consumers pause consumption during reassignment, (2) state/offset coordination takes time, (3) if a consumer left, remaining consumers suddenly get more partitions to handle.

**Example:** one consumer crashes → others inherit its partitions → lag jumps briefly.

## **27) Compare RabbitMQ vs Kafka in 3 points using the lecture’s “message queue tools” idea.**

**Answer:** RabbitMQ is a feature-rich broker for reliable queuing (AMQP), typically deletes after ack. Kafka is distributed streaming, optimized for high throughput and durable log storage with replay via offsets. Kafka fits event streams; RabbitMQ fits task queues.

**Example:** background email tasks → RabbitMQ; clickstream analytics → Kafka.

## 28) Explain “loose coupling” provided by message queues (3 points).

**Answer:** (1) Producer and consumer don't need to be online at the same time, (2) each service can scale independently, (3) failures don't immediately break the whole system because messages persist.

**Example:** payment service continues receiving orders even if invoice service is down temporarily.

## 29) Kafka's capabilities: list 3 and give an example for one.

**Answer:** Kafka can (1) publish/subscribe to streams, (2) store streams durably and fault-tolerantly, (3) process streams in real time via Kafka Streams/ksqlDB.

**Example:** real-time trending products computed from live clickstream.

## 30) What is Kafka Streams / ksqlDB / Kafka Connect? Mention 1 use case each.

**Answer:** **Kafka Connect** integrates Kafka with external systems (DB, S3, files). **Kafka Streams** is a Java library for stream processing on topics. **ksqlDB** provides SQL-like querying over streams.

**Example:** Connect pulls DB changes; Streams aggregates; ksqlDB queries live counts.

## 31) What is Schema Registry and why do we need it? (3 points)

**Answer:** Schema Registry stores message schemas (Avro/JSON/Protobuf) to ensure: (1) producer/consumer compatibility, (2) controlled schema evolution, (3) fewer runtime deserialization errors.

**Example:** adding a new optional field `discount` shouldn't break older consumers.

## **32) What does Zookeeper (legacy) do for Kafka? And what is KRaft?**

**Answer:** Zookeeper (legacy) manages cluster **metadata** and **controller elections**. KRaft is the newer approach replacing Zookeeper for metadata + controller management.

**Example:** when a controller fails, election selects a new one to manage partitions.

## **33) Scenario: Broker hosting a leader partition dies. Explain recovery in 4 steps.**

**Answer:** (1) Leader becomes unavailable, (2) Kafka selects a new leader from **in-sync replicas**, (3) producers/consumers redirect to new leader, (4) once failed broker returns, it may rejoin as follower and catch up.

**Example:** broker-1 down, broker-2 follower promoted to leader.

## **34) Explain “fault tolerance” in Kafka using replication (3 points).**

**Answer:** (1) Replicas stored across brokers protect against single broker failure, (2) leader/follower model enables automatic failover, (3) ISR ensures only fully caught-up replicas are trusted for durability.

**Example:** even if one broker dies, messages still available from replica leader.

## **35) Explain why partitions enable parallelism but still preserve ordering.**

**Answer:** Parallelism happens because different partitions can be processed by different consumers. Ordering is preserved because within a single partition, messages are strictly sequenced by offset. So you get “parallel across partitions, ordered within partition.”

**Example:** city=Karachi events in one partition ordered, Lahore in another partition processed in parallel.

## **36) Scenario: You need global ordering of all events. Can Kafka guarantee it? Explain properly.**

**Answer:** Kafka guarantees ordering **only within a partition**, not across partitions. To get global ordering you'd need: (1) single partition topic (hurts throughput), or (2) design keying so all relevant events go to one partition, or (3) external sequencing layer.

**Example:** single “transactions” partition ensures strict total order.

### **37) What should you monitor in Kafka and why? (at least 4 metrics + reason)**

**Answer:** Monitor: (1) broker health (availability), (2) consumer lag (freshness/backlog), (3) disk usage (Kafka stores logs), (4) network I/O (throughput), (5) GC activity (JVM pauses).

**Example:** rising lag + high GC  $\Rightarrow$  consumer slowdown due to memory pressure.

### **38) Name 3 monitoring tools from slides and what each is used for.**

**Answer:** (1) Prometheus + JMX Exporter: metrics collection from JVM/Kafka, (2) Grafana: dashboards, (3) Confluent Control Center: Kafka management/monitoring UI, (4) Burrow: consumer lag monitoring specifically.

**Example:** Burrow alerts when lag crosses threshold.

### **39) Scenario: Billing service and Analytics service both need every order event. Should they be in same consumer group?**

**Answer:** No. In one consumer group, each message is delivered to **only one** consumer in the group. To have both services get all messages: use **two separate consumer groups** (billing-group, analytics-group).

**Example:** billing-group consumes all orders; analytics-group also consumes all orders independently.

### **40) Scenario: You want to scale Billing service horizontally (multiple instances) but still process each order once. What Kafka feature enables this?**

**Answer:** Use **consumer group** for Billing: (1) multiple billing consumers in same group, (2) partitions distributed among them, (3) each message goes to only one billing consumer (prevents double processing inside same group).

**Example:** 3 partitions  $\Rightarrow$  up to 3 billing instances fully utilized.

### **41) Explain at-least-once delivery from the slide's data-flow. Where exactly does it come from?**

**Answer:** At-least-once occurs because: (1) Kafka stores events durably, (2) consumer keeps reading from leader and only advances progress by committing offset, (3) if consumer crashes before commit, it will re-read and reprocess => duplicates but no loss.

**Example:** crash after processing offset 50 but before committing ⇒ offset 50 processed again.

## 42) Scenario: Your consumer does heavy processing and falls behind. Give 3 fixes.

**Answer:** (1) Increase partitions + add consumers in group for parallelism, (2) optimize consumer work (batching, async processing, cheaper compute), (3) scale infra (more CPU/memory, better disk/network) and tune poll/commit strategy.

**Example:** move ML inference to async worker and only enqueue job id from consumer.

## 43) Explain the “Order Processing System” example using correct terms (producer, topic, consumer group).

**Answer:** Producer: Order Service sends **order events**. Topic: **orders** holds events. Consumer group: Billing Service + Analytics Service process orders (but correct design is separate groups if both need every message). Partitions enable parallelism; offsets track progress.

## 44) Numerical + reasoning: In a partition, latest offset=4 and consumer committed offset=2. What does lag=2 mean exactly? Name the messages pending.

**Answer:** Lag=2 means two messages after committed offset are waiting: offsets 3 and 4 are not yet processed/committed. It indicates backlog and potential delay.

**Example:** if offsets map to M4 and M5, those are pending.

## 45) “What happens if a consumer commits offset before processing is truly done?” (scenario, strict)

**Answer:** That risks **message loss** at application level: (1) consumer marks progress as done, (2) then crashes before real processing completes, (3) upon restart it starts after that offset, so that event is skipped by the app.

**Example:** committing before saving to DB can lose an order record.